

GPT Explained

A hands-on guide to transformer architecture

Ju Lin

Table of contents

Welcome	6
License	7
Preface	8
Who This Book is For?	8
How to Use This Book	8
Summary	9
Should I Buy This Book?	9
Contributors	10
 I Foundations	 11
1 Introduction	13
1.1 What Is GPT?	13
1.2 How GPT Works	14
1.3 A Brief History of GPT	16
 2 Notation and Definitions	 19
2.1 Scalars	19
2.2 Capital Sigma and Capital Pi	20
2.3 Vectors	21
2.4 Dot Product	25
2.5 Matrix Multiplication	29
2.6 Softmax	31
2.7 Logarithm and Exponential	33
2.8 Mean and Variance	37
2.9 Derivative	38
2.10 Key Takeaways	40

II	Representations	42
3	Tokens — Text to Numbers	44
3.1	The Idea	44
3.2	Byte Pair Encoding (BPE)	45
3.3	The Matrix: A Tiny Vocabulary Table	46
3.4	The Code: A Micro-BPE in Python	47
3.5	Key Takeaways	51
4	Embeddings — Numbers to Meaning	53
4.1	The Idea	53
4.2	Why Learned Vectors?	53
4.3	The Math	54
4.4	The Matrix: Worked Example	55
4.5	The Code: Embedding in Python	56
4.6	Key Takeaways	58
5	Positional Encoding — Giving Order to Meaning	61
5.1	The Idea	61
5.2	Sinusoidal Positional Encoding	62
5.3	The Math	62
5.4	The Matrix: Worked Example	63
5.5	The Code: Positional Encoding in Python	64
5.6	Learned vs Sinusoidal Positional Embeddings	67
5.7	Key Takeaways	67
III	Transformer Core	68
6	Attention — Tokens Talking to Each Other	70
6.1	The Idea	70
6.2	[source]	71
6.3	$\text{softmax}(z) = \exp(z) / \sum \exp(z)$	71
6.4	The Math	71
6.5	The Matrix: Worked Example	73
6.6	The Code: Scaled Dot-Product Attention in Python	74
6.7	Key Takeaways	77
7	Multi-Head Attention — Many Conversations at Once	78
7.1	The Idea	78
7.2	The Math	79
7.3	The Matrix: Worked Example	80
7.4	The Code: Multi-Head Attention in Python	81
7.5	Why Multi-Head Attention Works	82

7.6	Key Takeaways	83
8	Feed-Forward Network — The Model's Memory	84
8.1	The Idea	84
8.2	The Math	85
8.3	GELU Deep Dive	85
8.4	The Matrix: Worked Example	86
8.5	The Code: FFN in Python	87
8.6	The FFN as a Key-Value Memory	89
8.7	Key Takeaways	90
9	The Transformer Block — Putting It Together	91
9.1	The Idea	91
9.2	The Math	92
9.3	The Residual Stream	92
9.4	The Matrix: Worked Example	93
9.5	The Code: Transformer Block in Python	95
9.6	Why N Blocks?	97
9.7	Key Takeaways	97
IV	Prediction and Learning	99
10	Vocabulary Projection — From Vectors to Words	101
10.1	The Idea	101
10.2	The Math	101
10.3	The Generation Loop	102
10.4	The Matrix: Worked Example	103
10.5	The Code: Full microGPT Forward Pass in Python	105
10.6	Weight Tying in Detail	107
10.7	Key Takeaways	107
10.8	The Complete Picture	107
11	Loss — How the Model Learns	109
11.1	The Idea	109
11.2	The Problem: Evaluating a Probability Distribution	109
11.3	Cross-Entropy Loss	110
11.4	The Matrix: Worked Example	112
11.5	Python Implementation	113
11.6	Key Takeaways	115
12	Training — Teaching the Model	116
12.1	The Idea	116
12.2	The Goal: Move the Loss Down	116

12.3	The Chain Rule	117
12.4	Backprop Through Softmax and Cross-Entropy	118
12.5	Backprop Through a Linear Layer	118
12.6	Accumulating Gradients Across the Sequence	119
12.7	The Gradient Descent Step	119
12.8	One Training Step	120
12.9	Watching the Loss Fall	120
12.10	Key Takeaways	120
V	Modern Extensions	123
13	Modern GPT	125
13.1	Rotary Positional Encoding (RoPE)	125
13.2	Multi-Query and Grouped-Query Attention	127
13.3	Flash Attention	127
13.4	Alternative Architectures	128
13.5	Key Takeaways	129
	Appendices	130
A	microGPT in Python — Complete Runnable Code	130
A.1	Files	130
A.2	Installation & Running	130
A.3	End-to-End Demo	130
A.4	microGPT Architecture Summary (Reference)	131

Welcome

A hands-on guide to transformer architecture, from tokens to text generation, with math, matrices, Python, and diagrams.

This is the website for *GPT Explained*, an open book about how GPT-style language models actually work — from the inside out.

This book will teach you how transformers are built and trained. Not at a high level — at the level where you can read the code, follow the math, and trace a single token all the way through a real model.

You will learn tokenization, embeddings, positional encoding, attention, multi-head attention, feed-forward networks, transformer blocks, vocabulary projection, cross-entropy loss, and backpropagation. Each concept arrives as a focused chapter: a plain-language explanation, the math in full, a Python implementation you can run, and diagrams generated from code.

In this book, you will build a small but complete GPT from scratch. Not a toy metaphor — actual matrix multiplications, actual softmax, actual multi-head attention, actual gradient descent. By the final chapter you will have a working miniature language model you can inspect line by line, and a clear mental model of every step a modern LLM takes between reading your prompt and writing its reply. The only prerequisites are basic Python and enough math to follow an equation when it is explained in words next to it. No machine-learning background is assumed.

GPT Explained is free to read online under the [Creative Commons Attribution-NonCommercial-ShareAlike 3.0](#) license. You may share and adapt the material for non-commercial purposes as long as you give credit and keep the same license. Ebook and PDF versions are available for download on the releases page.

Contributions are welcome — whether that is a typo fix, a clearer explanation, a better diagram, or a new section. Open an issue or a pull request on [GitHub](#). Please read the [Code of Conduct](#) before contributing; we want this project to be a welcoming place for everyone learning how these models work.

License

This work is licensed under the Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License. To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc-sa/3.0/>.

Preface

Who This Book is For?

This book is for programmers who want to understand how GPT-style language models actually work — not just use them.

If you have written Python before and can follow basic algebra, you have enough background to read every page. You do not need prior experience with machine learning, deep learning, or neural networks. The book introduces every concept from scratch, explains the math in plain language before writing it as a formula, and implements each idea in code you can run and inspect.

If you are already a machine-learning practitioner, this book gives you a clean, self-contained reference for the transformer architecture. The implementations are deliberately simple — no frameworks, no abstractions beyond plain Python — so nothing hides what is actually happening.

How to Use This Book

Read the chapters in order on the first pass. Each chapter builds on the previous one: tokens before embeddings, embeddings before attention, attention before the full transformer block. Skipping ahead is possible but you may find yourself returning to fill gaps.

Every code snippet is part of a complete, runnable program. The source lives in `src/python/` and can be run with `python3 src/python/run_book_code.py`. Every diagram is generated from code in `src/matplotlib/`, `src/figures/`, or directly inside the chapter files. If a diagram surprises you, open its source and change the numbers.

Work through the exercises at the end of each chapter. The exercises are not optional decoration — they are the fastest way to verify that you understood the material rather than just read it.

Summary

The book covers thirteen chapters and one appendix:

Foundations. Chapter 1 gives an overview of GPT and its history. Chapter 2 introduces the math notation used throughout: vectors, matrices, dot products, softmax, logarithms, mean, and variance.

Representations. Chapter 3 explains tokenization and byte-pair encoding. Chapter 4 covers word embeddings and the embedding lookup. Chapter 5 introduces sinusoidal positional encoding.

Transformer Core. Chapter 6 derives scaled dot-product attention step by step. Chapter 7 extends it to multi-head attention. Chapter 8 covers the feed-forward sub-layer. Chapter 9 assembles these pieces into a full transformer block with residual connections and layer normalization.

Prediction and Learning. Chapter 10 covers vocabulary projection. Chapter 11 derives cross-entropy loss and perplexity. Chapter 12 traces backpropagation and the Adam optimizer through the full model.

Modern GPT. Chapter 13 describes the changes in current models: RoPE positional encoding, grouped-query attention, SwiGLU activations, and RMSNorm.

Appendix A presents the complete micro-GPT implementation — the same code from across the chapters assembled into one readable file.

Should I Buy This Book?

The full text is free to read online at the book’s website. Nothing is paywalled.

If you find the book useful and want to support the work, you can purchase the EPUB or PDF edition from the releases page. Buying a copy is a way of saying thank you — but it is entirely optional. The content is identical to what you are reading now.

Contributors

Contributors as of 377620c:

- Ju Lin

Part I

Foundations

This part establishes the vocabulary used throughout the rest of the book. It starts with the high-level idea of GPT-style models, then introduces the compact math notation needed to read the diagrams and Python code.

- In Chapter [1](#), you will see what a GPT model does and how the main pieces fit together.
- In Chapter [2](#), you will learn the vector, matrix, probability, and calculus notation used in the later chapters.

1 Introduction

This chapter answers two questions before you touch any math or code: what is GPT, and how does it work? We keep the language concrete and the diagrams simple. The details come later.

1.1 What Is GPT?

GPT stands for *Generative Pre-trained Transformer*.

Generative — it produces new text, one token at a time, rather than classifying or retrieving.

Pre-trained — before you ever use it, the model was trained on a vast corpus of text from the internet, books, and code. It learned the statistical patterns of language from billions of examples.

Transformer — the underlying neural network architecture. The transformer replaced recurrent networks in 2017 and became the backbone of every large language model since.

At its core, a GPT model is a function:

$$P(\text{next token} \mid \text{context tokens}) = f_{\theta}(\text{context})$$

Given a sequence of tokens (integers representing text), it outputs a probability distribution over the vocabulary — every possible next token gets a score. The model then samples from this distribution to generate the next word. Repeat, and you get fluent text.

Note

What is a token? A token is the smallest unit of text the model processes. It is not always a whole word — “unbelievable” might be two tokens: “unbel” and “ievable”. Chapter 3 explains exactly how text is split into tokens.

1.2 How GPT Works

Imagine texting with a really smart friend who has read every book, article, and website ever written. You send a message, they think for a split second, and they reply with exactly the right words. GPT works like that friend — except instead of a brain, it uses a chain of math operations called a *pipeline*.

Every word GPT produces goes through these seven steps, in order:

Step 1 — Tokenization (Chapter 3). Before GPT can think about your text, it has to chop it into small pieces called *tokens*. Think of tokens like LEGO bricks: you break a sentence apart, and GPT works with the individual bricks. The phrase "the cat" might become three bricks with serial numbers like [1, 428, 3797].

Step 2 — Embedding (Chapter 4). A number like 428 is just an ID — it has no meaning by itself. So GPT looks each ID up in a giant table and swaps it for a list of hundreds of numbers called a *vector*. Picture a map where every word has its own location. Similar words end up close together: "cat" and "kitten" are neighbors, while "cat" and "spaceship" are far apart.

Step 3 — Positional Encoding (Chapter 5). Here is a puzzle: does "dog bites man" mean the same thing as "man bites dog"? Obviously not — but if GPT only sees individual words and ignores their order, it cannot tell the difference. Positional encoding adds a tiny "position tag" to each word's vector so GPT knows which word came first, second, and so on.

Step 4 — Transformer Blocks (Chapters 6–9). This is where the real thinking happens. GPT passes the vectors through a stack of identical layers called *transformer blocks*. Inside each block, two things happen back-to-back:

- *Attention* — every word looks at every other word and asks "which of you matters most to me right now?" For example, the word "it" in "I kicked the ball and it bounced" needs to figure out that "it" refers to "ball".
- *Feed-forward* — each word's vector gets processed through a small math circuit that stores facts from training.

There can be dozens of these blocks stacked on top of each other. Each layer builds a slightly deeper understanding of the text.

Step 5 — Vocab Projection (Chapter 10). After passing through all the blocks, GPT has a rich, updated vector for the last word in your prompt. Now it needs to turn that into a real prediction: "what word should come next?" It does this by scoring every word in its vocabulary — all 50,000-plus of them. These raw scores are called *logits*.

Step 6 — Softmax (Chapter 10). Raw scores are hard to work with, so GPT converts them into proper probabilities that add up to 100%. The highest-probability word is the model's best guess. It samples from the top options to avoid always giving the same robotic answer.




Figure 1.1: The GPT pipeline: from input text to predicted next word, in seven steps

Step 7 — Loss and Training (Chapters 11–12). GPT did not start out smart. During training, it read billions of sentences and kept trying to predict the next word — getting it wrong most of the time at first. Each mistake triggered a feedback signal called *backpropagation* that nudged millions of internal numbers a tiny bit in the right direction. After trillions of these small nudges, the model got very good at the guessing game.

That’s GPT from start to finish. The rest of this book builds each step from scratch — in math, in Python code, and in step-by-step diagrams.

1.3 A Brief History of GPT

GPT did not appear out of nowhere. It is the result of more than sixty years of research — a long chain of ideas, failures, and breakthroughs. Here is the story, told simply.

1950s–1980s — The Birth of Machine Learning

Early computers followed strict rules written by programmers. If you wanted a computer to recognise a cat, you had to describe a cat in code: “four legs, pointy ears, fur...” That worked for simple things but fell apart for anything complicated.

Researchers started asking a different question: what if the computer could *learn the rules itself* by looking at lots of examples? That idea became *machine learning*. Instead of coding rules by hand, you feed the machine thousands of cat photos, label them “cat”, and let it figure out the pattern.

1943–1986 — Neural Networks

Inspired by the human brain, researchers built *neural networks* — programs made of many simple math nodes connected together, loosely like neurons. Each node takes some numbers in, does a small calculation, and passes the result on.

In 1986, a key paper by Geoffrey Hinton and others showed how to train these networks using an algorithm called *backpropagation*. Think of it like a teacher who marks your test, finds every mistake, and tells you exactly how much to adjust each answer. This was huge — but computers at the time were too slow for it to become widely useful yet.

2012 — The Deep Learning Explosion

Fast-forward to 2012. A neural network called AlexNet entered an image-recognition contest and demolished the competition. It won by a margin so wide that the whole field changed overnight. The secret ingredient was *deep learning* — stacking many layers of neurons on top of each other — combined with powerful graphics cards (GPUs) that could run the math fast enough.

Suddenly every company poured money into neural networks.

2013 — Word Vectors (Word2Vec)

A team at Google figured out how to teach a neural network to read text. They trained it to predict which words tend to appear near each other. The surprising result: the network assigned each word a list of numbers (a *vector*), and similar words ended up with similar numbers. You could even do word arithmetic: “king” – “man” + “woman” “queen”. Language had become math.

2014–2015 — Sequence Models and Attention

To translate a sentence — say, from English to French — models needed to process words in order. Researchers used networks called *RNNs* (Recurrent Neural Networks) and *LSTMs* for this. But they had a problem: long sentences broke them. By the time the model reached the end of a paragraph, it had mostly forgotten the beginning.

In 2015, researchers introduced *attention*. Instead of forcing the model to squeeze everything into one memory slot, attention let it look back at any earlier word at any time — like being allowed to flip back to any page of your notes during an exam. This made translation much better.

2017 — “Attention Is All You Need”

A team at Google published a paper with that bold title. They asked: what if we got rid of the old word-by-word processing entirely and used *only* attention? The result was the *transformer* — the architecture that GPT is built on. The transformer could look at all the words in a sentence at the same time, in parallel, which made it dramatically faster to train and better at long-range reasoning.

2018 — GPT-1 and BERT

Two important models appeared within months of each other.

OpenAI released *GPT-1*, which trained a transformer to do one simple task: read text and predict the next word. After seeing enough text, the model picked up grammar, facts, and reasoning as a side effect. It had 117 million *parameters* — the tunable numbers inside the network (imagine dials on a mixing board).

Google released *BERT*, which trained a transformer to predict missing words in the middle of a sentence instead of just the next one. BERT became the backbone of Google Search for years.

2019 — GPT-2

OpenAI scaled GPT-1 up to 1.5 billion parameters and fed it text scraped from the internet. The result was striking: GPT-2 could write paragraphs that sounded like a real person wrote them. OpenAI were nervous enough about misuse that they released it in stages rather than all at once.

2020 — GPT-3

Another scale jump: 175 billion parameters, trained on a huge chunk of the internet, books, and code. GPT-3 could write essays, answer questions, translate, summarise, and generate code — without being specifically trained for any of those things. It had just absorbed so much human writing that it learned to generalise. Developers started building products on top of it via an API, and a new industry was born.

2023 — LLaMA and the Open-Source Wave

Until now, the most powerful models were locked behind the walls of big companies. Meta changed that by releasing *LLaMA* — a family of models that could run on a single powerful laptop. Researchers and hobbyists could finally experiment with a state-of-the-art model without needing a data centre. Hundreds of variants appeared within weeks: Alpaca, Vicuna, Mistral, and many more. The open-source AI ecosystem exploded.

2022–2023 — ChatGPT and GPT-4

OpenAI added a training step called *reinforcement learning from human feedback* (RLHF). Human raters scored the model’s replies, and the model learned to be more helpful and less likely to say harmful things. The result was *ChatGPT*, released in November 2022 — the fastest app to reach 100 million users in history.

GPT-4 followed in 2023. It understood images as well as text, scored at or above human level on many standardised exams, and became the foundation for countless products and services.

Now — A Fast-Moving Field

Since then, new models have appeared every few months: Claude (Anthropic), Gemini (Google), Mistral, Qwen, DeepSeek, and many others. Each brings new ideas: longer context windows, better reasoning, cheaper inference, multilingual support.

The race is still on. But no matter how fancy the model, the fundamentals stay the same. Everything you learn in this book — tokens, embeddings, attention, transformer blocks, training — is the engine underneath all of them.

2 Notation and Definitions

This chapter is a compact reference for the core math used at the start of the book. It focuses on notation, definitions, and small numerical examples.

2.1 Scalars

A *scalar* is a single number. Examples include -2 , 0 , 0.5 , and 3.14 .

Scalars can be added, subtracted, multiplied, divided, squared, and compared. They are the simplest objects in this chapter. Vectors and matrices are built by arranging scalars into larger shapes.

Note

Math Minute – Variable

A variable is a name for a value. In $x = 3.0$, the variable is x and its value is 3.0 . In formulas, variables let us describe a whole pattern without writing a separate equation for every possible number.

2.1.1 Scalar Operations

For two scalars $a = 3.0$ and $b = 2.0$:

$$a + b = 5.0$$

$$ab = 6.0$$

$$a^2 = 9.0$$




Figure 2.1: Scalar operations combine individual numbers

2.2 Capital Sigma and Capital Pi

Capital sigma $\sum$ means “add a sequence of terms.” Capital pi $\prod$ means “multiply a sequence of terms.”

i Note

Math Minute – Index Variable

An index variable counts through positions in a sequence. In $\sum_{i=1}^4 x_i$, the index variable is i . It starts at 1, then takes the values 2, 3, and 4.

2.2.1 Capital Sigma: Sum

The expression:

$$\sum_{i=1}^4 x_i$$

means:

$$x_1 + x_2 + x_3 + x_4$$

For $x = [2, 3, 5, 7]$:

$$\sum_{i=1}^4 x_i = 2 + 3 + 5 + 7 = 17$$

i Note

Math Minute – Subscript

A subscript is the small label written below and to the right of a symbol. In x_i , the subscript i selects one entry from the sequence x . So x_1 means “the first entry of x .”

2.2.2 Capital Pi: Product

The expression:

$$\prod_{i=1}^4 x_i$$

means:

$$x_1 x_2 x_3 x_4$$

For $x = [2, 3, 5, 7]$:

$$\prod_{i=1}^4 x_i = 2 \cdot 3 \cdot 5 \cdot 7 = 210$$

2.3 Vectors

A *vector* $v \in \mathbb{R}^n$ is an ordered list of n real numbers: $v = [v_1, v_2, \dots, v_n]$. Vectors store many related measurements as one object. Once values are written as vectors, we can add, scale, compare, and transform them with ordinary arithmetic.




Figure 2.2: Capital sigma adds terms; capital pi multiplies terms

i Note

Math Minute — Element Of

The symbol $\in$ means “is an element of” or, more casually, “is in.” So $v \in \mathbb{R}^n$ says that v is one item from the set $\mathbb{R}^n$. In the above definition, v is a vector with n real-number components.

i Note

Math Minute — Real Number

A real number is any ordinary number on the number line: -2 , 0 , 0.5 , $\sqrt{2}$, π , and so on. The notation $\mathbb{R}$ means “the set of all real numbers.” So $\mathbb{R}^n$ means “all vectors with n real-number components.”

i Note

Math Minute — Vector Shape and Indexing

The shape of a vector is its length. If $v \in \mathbb{R}^n$, then v has shape n . The notation v_i means “entry i of vector v .” The notation v_n means “the last entry when the vector has length n .”

In Python, the book represents a vector as a list of floating-point numbers:

```
Vector = list[float]
```

This simple representation is enough for the small implementations in the book. Packages like NumPy provide their own vector type, usually as an `ndarray`, with faster arithmetic, compact memory layout, and many built-in operations. For larger numerical programs, you would usually use NumPy, PyTorch, JAX, or another tensor library instead.

v has shape 4




Figure 2.3: Vector shape and indexing

2.3.1 Vector Addition

When two vectors have the same length, we can add them by adding matching positions. The first component adds to the first component, the second to the second, and so on:

$$u + v = [u_1 + v_1, u_2 + v_2, \dots]$$

The Python version follows that sentence directly: walk through both vectors together and add each pair.

```
def add_vectors(left: Vector, right: Vector) -> Vector:
    return [a + b for a, b in zip(left, right)]
```

For example, take $u = [3.0, 4.0]$ and $v = [1.0, -2.0]$. The first components give $3.0 + 1.0 = 4.0$, and the second components give $4.0 + (-2.0) = 2.0$. So:

$$u + v = [3.0 + 1.0, 4.0 + (-2.0)] = [4.0, 2.0]$$

2.3.2 Scalar Multiplication

A scalar is a single number. Multiplying a vector by a scalar stretches or shrinks the whole vector without changing how many components it has:

$$c \cdot v = [cv_1, cv_2, \dots]$$

In code, this means multiplying every value in the list by the same scalar.

```
def scale_vector(vector: Vector, scalar: float) -> Vector:
    return [scalar * value for value in vector]
```

For $u = [3.0, 4.0]$, multiplying by 2.0 doubles both components:

$$2.0 \cdot u = [6.0, 8.0]$$

2.3.3 L2 Norm

The L2 norm is the length of a vector. For a two-dimensional vector, this is the Pythagorean theorem; for longer vectors, the same idea extends across all components:

$$\|v\| = \sqrt{v_1^2 + v_2^2 + \dots + v_n^2}$$

The implementation squares every component, adds the squares, then takes the square root.

```
def vector_norm(vector: Sequence[float]) -> float:
    return math.sqrt(sum(value * value for value in vector))
```

For $u = [3.0, 4.0]$, the length is:

$$\|u\| = \sqrt{3.0^2 + 4.0^2} = \sqrt{9.0 + 16.0} = 5.0$$

2.3.4 Unit Vector

A unit vector keeps the direction of the original vector but scales its length to 1. To get one, divide the vector by its own norm:

$$\hat{v} = v / \|v\|$$

The helper below combines the two earlier operations: compute the norm, then scale by its reciprocal.

```
def unit_vector(vector: Vector) -> Vector:
    return scale_vector(vector, 1.0 / vector_norm(vector))
```

For $u = [3.0, 4.0]$, the norm is 5.0, so each component is divided by 5.0:

$$[3.0/5.0, 4.0/5.0] = [0.6, 0.8]$$

Figure [2.4](#) shows vector addition and the L2 norm in 2D.

2.4 Dot Product

The *dot product* (also called inner product) combines two vectors into one scalar. It is the basic operation for asking, “how aligned are these two vectors?”

2.4.1 Dot Product

The dot product multiplies matching components and then adds the results:

$$u \cdot v = \sum_{i=1}^n u_i v_i = u_1 v_1 + u_2 v_2 + \dots + u_n v_n$$

That is also the most direct way to write it in Python. The `zip` pairs matching components, and `sum` adds the pairwise products.

```
def dot(left: Sequence[float], right: Sequence[float]) -> float:
    return sum(a * b for a, b in zip(left, right))
```




Figure 2.4: Vector addition and norms in 2D

With $u = [3.0, 4.0, 0.0]$ and $v = [1.0, -2.0, 5.0]$, the matching products are $3.0 \cdot 1.0$, $4.0 \cdot (-2.0)$, and $0.0 \cdot 5.0$. Adding them gives:

$$u \cdot v = (3.0)(1.0) + (4.0)(-2.0) + (0.0)(5.0) = 3.0 - 8.0 + 0.0 = -5.0$$

The sign matters. A positive dot product means the vectors point somewhat in the same direction. A negative dot product means they point somewhat against each other. If $u \cdot v = 0$, the vectors are orthogonal (perpendicular).

2.4.2 Cosine Similarity

The dot product mixes two things: vector length and direction. Cosine similarity removes the length part so we can compare direction alone. Starting from the geometric form:

$$u \cdot v = \|u\| \|v\| \cos \theta$$

we solve for $\cos \theta$:

$$\cos \theta = \frac{u \cdot v}{\|u\| \|v\|}$$

The implementation is just that formula: dot product divided by the two vector lengths.

```
def cosine_similarity(left: Sequence[float], right: Sequence[float]) -> float:
    return dot(left, right) / (vector_norm(left) * vector_norm(right))
```

Using the same u and v as above, $\|u\| = 5.0$ and $\|v\| = \sqrt{30.0}$. Since the dot product is -5.0 , the cosine similarity is:

$$\frac{u \cdot v}{\|u\| \|v\|} = \frac{-5.0}{5.0 \sqrt{30.0}} \approx -0.183$$

The value is below zero, so the vectors point in moderately opposite directions.

Figure [2.5](#) shows the dot product as a scalar projection in 2D.




Figure 2.5: Dot product as scalar projection in 2D

2.5 Matrix Multiplication

A matrix is a rectangular grid of numbers. Matrices make it possible to organize many vectors or many linear equations at once. Matrix multiplication combines rows from one matrix with columns from another. The book represents a matrix as a list of rows, where each row is a list of floats:

```
Matrix = list[list[float]]
```

Note

Math Minute – Matrix Shape and Indexing

A matrix with m rows and n columns has shape $m \times n$. The notation $A \in \mathbb{R}^{m \times n}$ means every entry of A is a real number, and A has m rows and n columns. The notation A_{ij} means “the entry in row i , column j of matrix A .”

A has shape 2×3

rows	1	2	3
	4	5	6
	columns		

$A[2,3] = 6$

Figure 2.6: Matrix shape and indexing

2.5.1 Transpose

Transposing a matrix swaps rows and columns. The first row becomes the first column, the second row becomes the second column, and so on. If $A \in \mathbb{R}^{m \times n}$, then $A^T \in \mathbb{R}^{n \times m}$.

In Python, `zip(*matrix)` groups together the first item from every row, then the second item from every row, and so on. Converting each group back to a list gives the transposed rows.

```
def transpose(matrix: Matrix) -> Matrix:
    return [list(col) for col in zip(*matrix)]
```

For:

$$A = \begin{bmatrix} 1.0 & 2.0 \\ 3.0 & 4.0 \end{bmatrix}$$

the transpose is:

$$A^T = \begin{bmatrix} 1.0 & 3.0 \\ 2.0 & 4.0 \end{bmatrix}$$

2.5.2 Matrix Multiplication

Matrix multiplication combines rows from the left matrix with columns from the right matrix. Given $A \in \mathbb{R}^{m \times k}$ and $B \in \mathbb{R}^{k \times n}$, their product $C = AB$ has shape $m \times n$, and each entry is:

$$C_{ij} = \sum_{r=1}^k A_{ir} B_{rj}$$

That equation says: take row i from A , take column j from B , and compute their dot product. The implementation transposes the right matrix first so its columns become easy to iterate over as rows.

```
def matrix_multiply(left: Matrix, right: Matrix) -> Matrix:
    right_t = transpose(right)
    return [[dot(row, col) for col in right_t] for row in left]
```

For:

$$A = \begin{bmatrix} 1.0 & 2.0 \\ 3.0 & 4.0 \end{bmatrix}$$

and:

$$B = \begin{bmatrix} 5.0 & 6.0 \\ 7.0 & 8.0 \end{bmatrix}$$

the four output cells are:

$$C_{00} = 1.0 \cdot 5.0 + 2.0 \cdot 7.0 = 19.0$$

$$C_{01} = 1.0 \cdot 6.0 + 2.0 \cdot 8.0 = 22.0$$

$$C_{10} = 3.0 \cdot 5.0 + 4.0 \cdot 7.0 = 43.0$$

$$C_{11} = 3.0 \cdot 6.0 + 4.0 \cdot 8.0 = 50.0$$

So:

$$AB = \begin{bmatrix} 19.0 & 22.0 \\ 43.0 & 50.0 \end{bmatrix}$$

Figure 2.7 illustrates how each output cell is a dot product of a row with a column.

2.6 Softmax

The *softmax* function turns arbitrary real numbers into probabilities. The input numbers are called logits. The output values are all positive and sum to 1.

2.6.1 Vector Softmax

Softmax exponentiates each logit and divides by the total exponentiated mass:

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

In code, we subtract the maximum logit before exponentiating. This does not change the final probabilities, because the same constant shift appears in the numerator and denominator. It does prevent very large exponentials from overflowing.

1	2		5	6	19	22
3	4		7	8	43	50
A		X	B		=	C

$$C[1,1] = 1 \times 5 + 2 \times 7 = 19$$

Figure 2.7: Matrix multiplication: each output cell is a dot product of a row with a column

$$\text{softmax}(z)_i = \frac{e^{z_i - \max(z)}}{\sum_j e^{z_j - \max(z)}}$$

```
def softmax(logits: Sequence[float]) -> Vector:
    max_logit = max(logits)
    exp_values = [math.exp(value - max_logit) for value in logits]
    total = sum(exp_values)
    return [value / total for value in exp_values]
```

For $z = [2.0, 1.0, 0.1]$, the maximum is 2.0, so the shifted values are $[0.0, -1.0, -1.9]$. After exponentiating, we get approximately $[1.000, 0.368, 0.150]$, whose sum is 1.518. Dividing each value by that sum gives:

$$\text{softmax}(z) \approx [0.659, 0.242, 0.099]$$

The largest logit gets the largest probability, but the smaller logits still receive nonzero probability.

2.6.2 Row-wise Softmax

Sometimes the input is a matrix of scores. Each row can be normalized into its own probability distribution. Row-wise softmax applies the vector softmax independently to every row.

```
def softmax_rows(matrix: Matrix) -> Matrix:
    return [softmax(row) for row in matrix]
```

After this step, each row sums to 1.

Figure 2.8 shows how softmax maps raw logits to a probability distribution.

2.7 Logarithm and Exponential

The natural exponential e^x and natural logarithm $\ln x$ are inverse operations:

$$e^{\ln x} = x \quad \text{and} \quad \ln(e^x) = x$$

These functions convert between additive and multiplicative scales. They are useful whenever growth, decay, ratios, or repeated multiplication need to be written on an additive scale.




Figure 2.8: Softmax turns logits into probabilities — higher values dominate, none reach zero

2.7.1 Natural Exponential

The natural exponential raises e to a power:

$$\exp(x) = e^x$$

In Python, this is `math.exp`.

```
def natural_exp(value: float) -> float:
    return math.exp(value)
```

Two anchor values are worth remembering. If the exponent is 0.0, the result is 1.0. If the exponent is 1.0, the result is e :

$$e^{0.0} = 1.0$$

$$e^{1.0} \approx 2.718$$

2.7.2 Natural Logarithm

The natural logarithm asks the reverse question: what power of e gives this value?

$$\ln x = y \quad \text{means} \quad e^y = x$$

In Python, this is `math.log`.

```
def natural_log(value: float) -> float:
    return math.log(value)
```

Some useful values:

$$\ln(1.0) = 0.0$$

$$\ln(e) = 1.0$$

$$\ln(0.5) \approx -0.693$$

$$\ln(0.1) \approx -2.303$$

For inputs between 0 and 1, the natural logarithm is negative. Negating it gives a positive value that grows as the input moves closer to zero, as shown in Figure 2.9. For example:

$$-\ln(0.01) = 4.605$$




Figure 2.9: The natural log and its negative over inputs between 0 and 1

Two identities appear often when simplifying logarithms:

$$\ln(ab) = \ln a + \ln b$$

$$\ln(a^n) = n \ln a$$

2.8 Mean and Variance

Mean and variance summarize a vector of numbers. They tell us where the values are centered and how spread out they are.

2.8.1 Mean

The mean is the average value. Add all values, then divide by how many values there are:

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i$$

The implementation is just that calculation.

```
def mean(values: Sequence[float]) -> float:
    return sum(values) / len(values)
```

For $x = [2.0, 4.0, 4.0, 4.0, 5.0, 5.0, 7.0, 9.0]$, the sum is 40.0 and there are 8 values. So:

$$\mu = \frac{40.0}{8.0} = 5.0$$

2.8.2 Variance

Variance measures how spread out the values are. First subtract the mean from each value. Then square those deviations so negative and positive deviations do not cancel out. Finally, average the squared deviations:

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

The Python helper computes the mean first, then follows the formula.

```
def variance(values: Sequence[float]) -> float:
    avg = mean(values)
    return sum((value - avg) ** 2 for value in values) / len(values)
```

Continuing with $x = [2.0, 4.0, 4.0, 4.0, 5.0, 5.0, 7.0, 9.0]$, the mean is 5.0. The squared deviations are:

9.0, 1.0, 1.0, 1.0, 0.0, 0.0, 4.0, 16.0

Their average is:

$$\sigma^2 = \frac{9.0 + 1.0 + 1.0 + 1.0 + 0.0 + 0.0 + 4.0 + 16.0}{8.0} = 4.0$$

2.8.3 Standard Deviation

Standard deviation puts variance back on the original scale by taking the square root:

$$\sigma = \sqrt{\sigma^2}$$

```
def standard_deviation(values: Sequence[float]) -> float:
    return math.sqrt(variance(values))
```

In this example, the variance is 4.0, so the standard deviation is:

$$\sigma = \sqrt{4.0} = 2.0$$

Figure [2.10](#) illustrates the example data with mean and standard deviation marked.

2.9 Derivative

A *derivative* measures how fast one value changes as another value changes. For a function $f(x)$, the derivative is written:

$$\frac{df}{dx}$$

This reads as “the derivative of f with respect to x .” It tells us the slope of the function at a particular input.




Figure 2.10: Data points with mean and standard deviation

i Note

Math Minute – Function

A function maps an input to an output. If $f(x) = x^2$, then $f(3) = 9$. The input is x and the output is $f(x)$.

i Note

Math Minute – Slope

Slope means “change in vertical value divided by change in horizontal value.” Between two points, slope is:

$$\frac{\Delta y}{\Delta x}$$

A derivative is the slope when the two points get infinitely close together.

2.9.1 Example

For:

$$f(x) = x^2$$

the derivative is:

$$\frac{df}{dx} = 2x$$

At $x = 3$, the slope is:

$$2x = 2 \cdot 3 = 6$$




Figure 2.11: A derivative is the local slope of a function

2.10 Key Takeaways

This chapter defined a compact mathematical toolkit.

Concept	Main idea
Scalar	A single number
Capital Sigma/Pi	Compact notation for repeated addition or multiplication
Vector	An ordered list of numbers

Concept	Main idea
Dot product	A scalar measure of alignment between two vectors
Matrix multiply	Row-by-column dot products arranged into a new matrix
Softmax	A way to convert real-valued scores into probabilities
Logarithm	The inverse of exponentiation
Mean/Variance	Measurements of center and spread
Derivative	The local rate of change of a function

Together, these tools cover the notation and arithmetic used throughout the examples in this chapter.

Part II

Representations

This part turns raw text into vectors the model can compute with. The goal is to make each row of the model’s input matrix carry both identity and position.

- In Chapter 3, you will see how text becomes token IDs, and how a small byte-pair encoding tokenizer works.
- In Chapter 4, token IDs become learned vectors through an embedding matrix.
- In Chapter 5, those vectors gain order so the model can distinguish “dog bites man” from “man bites dog.”

3 Tokens — Text to Numbers

A language model does not read text. It reads *integers*. The first job of any GPT pipeline is to convert a string of characters into a sequence of integers. That conversion is called *tokenization*, and the integers it produces are called *token IDs*.

This chapter explains what tokens are, how they are produced, and what they look like as data structures. We implement a tiny byte-pair encoding tokenizer in Python.

3.1 The Idea

Consider the sentence:

```
the cat sat on the mat
```

A human reads six words. A GPT model reads something like:

```
[1the] [2cat] [3sat] [4on] [1the] [5mat]
```

Where each bracketed item is a token and the number is its ID in a vocabulary table. The model never sees the string "the" — it only ever sees the integer 1.

Why not just use characters? You could map each letter to an integer: a→1, b→2, The problem is that the model would need to learn that c, a, t together mean something different from c, a, r. It would have to rediscover every word from scratch. This makes learning extremely sample-inefficient.

Why not use whole words? The opposite extreme. Map every word to an integer. This works for common words, but English has hundreds of thousands of words, and new ones appear constantly. The vocabulary would be enormous, and rare words would have almost no training data.

Subword tokenization finds the middle ground. Common whole words become single tokens. Rare words get split into recognizable pieces. The word "tokenization" might become ["token", "ization"]. The word "running" might become ["run", "ning"].

i Note

Math Minute — Sets and Mappings

A vocabulary V is a finite set of strings. A tokenizer is a function $T : \text{String} \rightarrow \text{List}[\mathbb{Z}]$ mapping a string to a sequence of integers from $\{0, 1, \dots, |V| - 1\}$. Each integer is an index into V .

3.2 Byte Pair Encoding (BPE)

The most common tokenization algorithm used in GPT models is *Byte Pair Encoding*. The algorithm has two phases:

Training phase (happens once, offline):

1. Start with a vocabulary of individual characters (or bytes).
2. Count every adjacent pair of tokens in the training corpus.
3. Merge the most frequent pair into a single new token.
4. Repeat until the vocabulary reaches the target size (e.g., 50,000).

Encoding phase (happens at inference time):

1. Start with the input string split into individual characters.
2. Greedily merge adjacent pairs according to the trained merge rules, highest-priority first.
3. Return the final token IDs.

3.2.1 The Math

Let $V_0 = \{c_1, c_2, \dots\}$ be the initial character vocabulary. After each merge step k :

$$V_{k+1} = V_k \cup \{(a, b) \mid \text{freq}(a, b) \text{ is max over all adjacent pairs}\}$$

The merge order defines a priority list $M = [(a_1, b_1), (a_2, b_2), \dots]$. To encode a string:

```
tokens = characters(string)
for (a, b) in M:
    tokens = apply_merge(tokens, a, b)
return [vocab_id(t) for t in tokens]
```

Where `apply_merge` replaces every adjacent occurrence of `(a, b)` with the merged token `ab`.

3.3 The Matrix: A Tiny Vocabulary Table

Let's trace through a micro-example. Suppose our entire training corpus is:

```
low lower newest
```

Step 0 — Character vocabulary:

ID	Token
0	l
1	o
2	w
3	e
4	r
5	n
6	s
7	t
8	_ ← end-of-word marker

Initial tokenized corpus: [l,o,w,_] [l,o,w,e,r,_] [n,e,w,e,s,t,_]

Step 1 — Count pairs:

```
(l,o): 2    ← most frequent!  
(o,w): 2  
(w,_): 1  
(w,e): 1  
...
```

Merge (l,o) → lo. Vocabulary now has 9 entries: lo gets ID 9.


Step 2 — Corpus becomes:

```
[lo,w,_] [lo,w,e,r,_] [n,e,w,e,s,t,_]
```

Most frequent pair now: (lo,w) with count 2. Merge → low (ID 10).

After a few more merges, encoding "low" returns [10] — a single token.

The following figure shows the BPE merge sequence used to turn frequent character pairs into



longer tokens.

3.4 The Code: A Micro-BPE in Python

Python is used throughout this book — it strips away syntactic ceremony so the *structure of computation* is visible. We build from named abstractions upward, one chunk at a time.

The vocabulary abstraction. Before writing any BPE logic, we name what a vocabulary *is*. A vocab maps string tokens to integer IDs. We expose three operations — create, extend, lookup — and never touch the underlying list structure from outside:

```
def make_vocab() -> dict[str, int]:
    return {}

def vocab_extend(vocab: dict[str, int], token: str, token_id: int) -> dict[str, int]:
    return vocab | {token: token_id}

def vocab_lookup(vocab: dict[str, int], token: str) -> int | None:
    return vocab.get(token)
```

The idea: name what you need, implement it once, and never revisit the representation again.

Characters and corpus. A word enters BPE as a list of single-character strings, followed by "_" as an end-of-word marker. `words`->`corpus` is simply `map` over words.

```
def word_to_tokens(word: str) -> list[str]:
    return list(word) + ["_"]

def words_to_corpus(words: Sequence[str]) -> list[list[str]]:
    return [word_to_tokens(word) for word in words]
```

Counting pairs. Every adjacent pair (a b) in the corpus is a merge candidate.

```
def adjacent_pairs(tokens: Sequence[str]) -> list[tuple[str, str]]:
    return list(zip(tokens, tokens[1:]))

def count_pairs(corpus: Sequence[Sequence[str]]) -> Counter[tuple[str, str]]:
    return Counter(pair for tokens in corpus for pair in adjacent_pairs(tokens))

def most_frequent(counts: Counter[tuple[str, str]]) -> tuple[str, str]:
    return max(counts, key=counts.get)
```

Applying a merge. Replacing every (a b) pair in a token sequence with the concatenated string "ab" has a natural recursive structure.

```
def merge_pair(tokens: Sequence[str], left: str, right: str) -> list[str]:
    merged = []
    i = 0
    while i < len(tokens):
        if i + 1 < len(tokens) and tokens[i] == left and tokens[i + 1] == right:
            merged.append(left + right)
            i += 2
        else:
            merged.append(tokens[i])
            i += 1
    return merged

def merge_corpus(corpus: Sequence[Sequence[str]], left: str, right: str) -> list[list[str]]:
    return [merge_pair(tokens, left, right) for tokens in corpus]
```


The training loop. `train-bpe` is a single named-let with three accumulators: the current corpus, the growing vocabulary, and the ordered list of merges.

```
def unique_tokens(corpus: Sequence[Sequence[str]]) -> list[str]:
    return sorted({token for tokens in corpus for token in tokens})

def train_bpe(words: Sequence[str], target_size: int) -> tuple[dict[str, int], list[tuple[str, str]]]:
    corpus = words_to_corpus(words)
    vocab = {token: token_id for token_id, token in enumerate(unique_tokens(corpus))}
    merges = []

    while len(vocab) < target_size:
        counts = count_pairs(corpus)
        if not counts:
            break
        left, right = most_frequent(counts)
        new_token = left + right
        if new_token not in vocab:
            vocab[new_token] = len(vocab)
        merges.append((left, right))
        corpus = merge_corpus(corpus, left, right)

    return vocab, merges
```

Encoding. Given trained merge rules, encoding applies them in order to a fresh character sequence.

```
def encode(word: str, merges: Sequence[tuple[str, str]], vocab: dict[str, int]) -> list[int]:
    tokens = word_to_tokens(word)
    for left, right in merges:
        tokens = merge_pair(tokens, left, right)
    missing = [token for token in tokens if token not in vocab]
    if missing:
        raise ValueError(f"unknown token(s): {missing}")
    return [vocab[token] for token in tokens]
```

Demo. Run this with `python3 src/python/chapter_demos.py`:

```
def demo() -> dict[str, object]:
    words = ["low", "lower", "newest", "widest"]
    vocab, merges = train_bpe(words, 20)
```

```

return {
    "vocab": vocab,
    "merges": merges,
    "low": encode("low", merges, vocab),
    "lower": encode("lower", merges, vocab),
}

```

Expected output:

Vocabulary:

```


0 -> 1
1 -> o
...
9 -> lo
10 -> low
...

```

Encoding "low": (10)


Encoding "lower": (10 3 4 8)

The following figure shows how each BPE step replaces the most frequent byte pair with a new to-



ken.

The following figure shows the full tokenization pipeline from raw text to integer IDs.



3.5 Key Takeaways

- Text enters the model as a *sequence of integers* (token IDs).
- Tokens are *subword units*, not characters or whole words.
- The *vocabulary* is the fixed lookup table mapping tokens to IDs.
- BPE builds the vocabulary by *greedily merging* the most frequent pairs.
- The tokenizer is trained *once*; at inference time it only applies the stored merge rules.

i Note

What's next? We have a sequence of integers: [10, 3, 4, 8, ...]. These integers are just row indices. To give them geometric meaning — so the model can do math on them — we need an *embedding matrix*. That's Chapter 2.

4 Embeddings — Numbers to Meaning

We have token IDs: [10, 3, 4, 8]. They are integers. The model cannot do useful math on raw integers — adding $10 + 3$ gives 13, which means nothing. To do useful math, each integer needs to become a *vector* — a list of numbers that encodes the token’s meaning geometrically.

The mechanism is called *the embedding matrix*, and it is the model’s first learned representation.

4.1 The Idea

A token ID is just a number — say, 42. That integer is an address, not a meaning. You cannot do useful arithmetic on addresses: $42 + 7 = 49$ tells you nothing about the relationship between the two words.

To give each token a meaning the model can actually compute with, we replace its ID with a *list of hundreds of decimal numbers* — a vector. Think of it like this: the model keeps a big table, one row per word in the vocabulary. When it sees token 42, it pulls out row 42 and uses those numbers as the token’s representation going forward.

The key property that makes this useful is that the table is *learned*. After training on billions of words, tokens with similar meanings end up with similar rows. “Cat” and “kitten” cluster together. “Run” and “sprint” are nearby. “Cat” and “justice” are far apart. The model discovers these relationships entirely from how words appear together in text — no human labels required.

This learned table is the *embedding matrix*, and the step that converts a sequence of IDs into a sequence of vectors is called *embedding*.

4.2 Why Learned Vectors?

The values in the embedding matrix are not hand-crafted. They are initialized randomly and then *learned via gradient descent* during training, just like any other model parameter.

The training objective is next-token prediction: given $t_1, t_2, \dots, t_n$, predict t_{n+1} . Because the model must do this well across billions of examples, it is forced to arrange embeddings such that tokens that appear in similar contexts land near each other in the embedding space.

This produces the famous arithmetic property:

$$\text{embedding}(\text{king}) - \text{embedding}(\text{man}) + \text{embedding}(\text{woman}) \approx \text{embedding}(\text{queen})$$

No one programmed this. It is a geometric consequence of how the vectors were shaped by the training signal.

i Note

Math Minute — Dot Product

The dot product of two vectors $u = [u_1, \dots, u_n]$ and $v = [v_1, \dots, v_n]$ is: $u \cdot v = u_1v_1 + u_2v_2 + \dots + u_nv_n$. It is a single number. If both vectors point in the same direction, the dot product is large and positive. If they are perpendicular, it is zero. If they point opposite, it is negative. Dot product measures *alignment*.

4.3 The Math

Given:

- $|V|$ = vocabulary size
- $d = d_{\text{model}}$ (hidden dimension)
- $E \in \mathbb{R}^{|V| \times d}$ = embedding matrix (learned)
- $i \in \{0, \dots, |V| - 1\}$ = token ID

The embedding of token i is:

$$e_i = E_{i,:} \in \mathbb{R}^d$$

(Row i of matrix E .)

For a sequence of T tokens $[i_1, \dots, i_T]$, the embedding operation is:

$$X = E[[i_1, \dots, i_T], :] \in \mathbb{R}^{T \times d}$$

This is just row indexing — no multiplication. In practice it is implemented as a matrix multiplication with a one-hot vector:

$$e_i = o_i^\top \cdot E, \quad \text{where } o_i \in \{0, 1\}^{|V|}, \quad (o_i)_j = \begin{cases} 1 & j = i \\ 0 & j \neq i \end{cases}$$

But implementations use the index directly because it is faster.

4.4 The Matrix: Worked Example

Let's use tiny numbers: $|V| = 5$, $d = 4$.

Embedding matrix E (5×4):

```
      col0  col1  col2  col3
tok 0: [ 0.10 -0.20  0.30 -0.40 ]
tok 1: [ 0.50  0.60 -0.70  0.80 ]
tok 2: [-0.90  0.10  0.20 -0.30 ]
tok 3: [ 0.40 -0.50  0.60 -0.70 ]
tok 4: [-0.10  0.80 -0.40  0.50 ]
```

Input token sequence: "low lower" $\rightarrow$ token IDs [2, 3, 0] (hypothetical).

Embedding lookup:

```
X = E[[2, 3, 0], :]  
  
X[0] = E[2] = [-0.90,  0.10,  0.20, -0.30]   $\leftarrow$  "low"  
X[1] = E[3] = [ 0.40, -0.50,  0.60, -0.70]   $\leftarrow$  "low" (part of "lower")  
X[2] = E[0] = [ 0.10, -0.20,  0.30, -0.40]   $\leftarrow$  "er"  
  
Result X: shape [3  $\times$  4]
```

This $[3 \times 4]$ matrix is what flows into the next stage.

Figure [Figure 4.1](#) shows token IDs selecting row vectors from the embedding matrix.

Semantic similarity in vector space:

If we compute the dot product of two token embeddings, we get a scalar measuring how aligned they are:

```
dot(E[0], E[2]) = (0.10)(-0.90) + (-0.20)(0.10) + (0.30)(0.20) + (-0.40)(-0.30)
                 = -0.09 - 0.02 + 0.06 + 0.12
                 = 0.07   (slightly positive  $\rightarrow$  slight similarity)
```

After training, semantically related tokens will have much higher dot products.




Figure 4.1: Token IDs mapped to embedding rows

4.5 The Code: Embedding in Python

Before writing a single line of embedding logic, we name what a matrix *is*. A matrix is a rows-by-cols grid of numbers. We store it as a flat vector in row-major order, but nothing above this interface will ever need to know that.

```
def make_matrix(rows: int, cols: int, fill: float = 0.0) -> Matrix:
    return [[fill for _ in range(cols)] for _ in range(rows)]

def shape(matrix: Matrix) -> tuple[int, int]:
    return len(matrix), len(matrix[0]) if matrix else 0
```

`random_matrix` uses Xavier initialization: it scales the random values by $\sqrt{2/(fan_{in} + fan_{out})}$, which keeps activations from exploding or vanishing.

```
def random_matrix(
    rows: int,
    cols: int,
    rng: random.Random,
    scale: float | None = None,
) -> Matrix:
```



```
scale = math.sqrt(2.0 / (rows + cols)) if scale is None else scale
return [[rng.uniform(-scale, scale) for _ in range(cols)] for _ in range(rows)]
```

The embedding of token i is literally row i of the matrix — a table lookup.

```
def embed_token(embedding: Matrix, token_id: int) -> Vector:
    return list(embedding[token_id])
```

A sequence of T token IDs calls for T row lookups stacked into a $[T \times d]$ matrix.

```
def embed_sequence(embedding: Matrix, token_ids: Sequence[int]) -> Matrix:
    return [embed_token(embedding, token_id) for token_id in token_ids]
```

The dot product $u \cdot v = \sum_i u_i v_i$ measures geometric alignment.

```
# tag::vector_ops[]
# tag::add_vectors[]
def add_vectors(left: Vector, right: Vector) -> Vector:
    return [a + b for a, b in zip(left, right)]
# end::add_vectors[]

# tag::scale_vector[]
def scale_vector(vector: Vector, scalar: float) -> Vector:
    return [scalar * value for value in vector]
# end::scale_vector[]

# tag::vector_norm[]
def vector_norm(vector: Sequence[float]) -> float:
    return math.sqrt(sum(value * value for value in vector))
# end::vector_norm[]

# tag::unit_vector[]
def unit_vector(vector: Vector) -> Vector:
    return scale_vector(vector, 1.0 / vector_norm(vector))
# end::unit_vector[]
# end::vector_ops[]
```

```
# tag::dot[]
def dot(left: Sequence[float], right: Sequence[float]) -> float:
    return sum(a * b for a, b in zip(left, right))
# end::dot[]

# tag::cosine_similarity[]
def cosine_similarity(left: Sequence[float], right: Sequence[float]) -> float:
    return dot(left, right) / (vector_norm(left) * vector_norm(right))
# end::cosine_similarity[]
```

Demo. Run with `python3 src/python/chapter_demos.py`:

```
def chapter_04(seed: int = 4) -> dict[str, object]:
    rng = random.Random(seed)
    embedding = random_matrix(100, 8, rng)
    x = embed_sequence(embedding, [1, 2, 3])
    return {
        "sequence_shape": (len(x), len(x[0])),
        "dot": dot(embedding[1], embedding[2]),
        "cosine": cosine_similarity(embedding[1], embedding[2]),
    }
```

Figure Figure 4.2 shows how each token ID indexes into one row of the embedding weight matrix.

Figure Figure 4.3 shows how semantic relationships can appear as vector directions in embedding space.

4.6 Key Takeaways

- An embedding matrix $E \in \mathbb{R}^{|V| \times d}$ maps token IDs to dense vectors.
- The embedding of token i is row i of E — a single *lookup*, no computation.
- For a sequence of T tokens, we get a $[T \times d]$ matrix X .
- The vectors are *learned*: training nudges similar tokens closer together.
- Semantic arithmetic ($\text{king} - \text{man} + \text{woman} \approx \text{queen}$) emerges from the training objective.




Figure 4.2: Embedding matrix lookup

i Note

What's next? We have a matrix $\mathbf{X}$ of shape $[T \times d]$. Each row knows *what* the token is, but nothing about *where* it sits in the sequence. Swapping two tokens would produce the same rows in a different order, which the model would treat identically. We need to inject *position information* — that's Chapter 3.




Figure 4.3: Semantic arithmetic in embedding space

5 Positional Encoding — Giving Order to Meaning

After the embedding step we have a matrix $X \in \mathbb{R}^{T \times d}$. Each row $X[t]$ is a vector representing token t . But here is the problem: attention processes all tokens *simultaneously* — there is no inherent notion of “token 3 comes after token 2.” If you shuffled the rows, the model would not know.

We fix this by *adding a position vector* to each token’s embedding. The combined vector carries both *what* (the token identity) and *where* (the position).

5.1 The Idea

After embedding, every token is a vector. But those vectors carry no information about *where* in the sentence the token appears. “Dog bites man” and “man bites dog” produce the same three vectors, just in a different order — and attention would treat both sentences identically if we did nothing.

The fix is straightforward: before passing the vectors into attention, we *add a position signal* to each one. Token 0 gets a small nudge in one direction; token 1 gets a different nudge; token 50 gets yet another. The nudge is itself a vector the same size as the embedding, and it is designed so that no two positions produce the same nudge.

Think of it like stamping each page of a manuscript with its page number before sending it to the printer. The text on the page does not change — but now the order is recoverable.

There are two ways to design the position stamp:

- *Learned*: keep a second lookup table, one row per position, and let the model learn the best stamps from data. GPT-2 does this.
- *Sinusoidal*: compute the stamp from a fixed mathematical formula using waves of different frequencies. The original transformer paper did this. We study both, because the sinusoidal approach illuminates *why* position signals work.

5.2 Sinusoidal Positional Encoding

For position $\mathbf{t}$ and dimension i (where $0 \leq i < d$):

$$PE(t, 2i) = \sin\left(\frac{t}{10000^{2i/d}}\right)$$
$$PE(t, 2i + 1) = \cos\left(\frac{t}{10000^{2i/d}}\right)$$

Even-indexed dimensions use sine; odd-indexed dimensions use cosine. The denominator $10000^{2i/d}$ controls the frequency — it grows exponentially with i , so low dimensions have high frequency (vary rapidly with $\mathbf{t}$) and high dimensions have low frequency (vary slowly).

Why this formula works:

1. *Unique per position:* Every $\mathbf{t}$ produces a distinct vector.
2. *Bounded:* All values lie in $[-1, 1]$, so they do not dominate the token embeddings.
3. *Smooth:* Nearby positions have similar encodings, so the model can learn “close in position $\rightarrow$ related in meaning.”
4. *Extrapolatable:* The formula works for any $\mathbf{t}$, even positions not seen during training.

A key property: The encoding for position $\mathbf{t} + \mathbf{k}$ can be expressed as a *linear transformation* of the encoding for position $\mathbf{t}$. This means the model can learn to attend to “the token 3 positions back” using a simple linear operation — no memorization of absolute positions required.

5.3 The Math

Full formula for the positional encoding matrix $PE \in \mathbb{R}^{T \times d}$:

$$PE[t, i] = \sin(t \cdot \omega_{i/2}) \quad \text{if } i \text{ is even}$$
$$PE[t, i] = \cos(t \cdot \omega_{\lfloor i/2 \rfloor}) \quad \text{if } i \text{ is odd}$$

where $\omega_k = 1/10000^{2k/d}$.

After building PE, the position-encoded input is:

$$\tilde{X} = X + PE \in \mathbb{R}^{T \times d}$$

The model now works with $\tilde{X}$ instead of X .

5.4 The Matrix: Worked Example

Let $T = 4$ (4 tokens), $d = 8$ (8-dimensional).

First compute the frequencies ω_k for $k = 0, 1, 2, 3$:

$$\begin{aligned}\omega_0 &= 1/10000^{0/8} = 1.0 \\ \omega_1 &= 1/10000^{2/8} \approx 0.1 \\ \omega_2 &= 1/10000^{4/8} = 0.01 \\ \omega_3 &= 1/10000^{6/8} \approx 0.001\end{aligned}$$

Now compute PE row by row (position $t = 0, 1, 2, 3$):

```
PE[0] = [sin(0·1.0), cos(0·1.0), sin(0·0.1), cos(0·0.1), ...]
       = [0.000,      1.000,      0.000,      1.000,      0.000, 1.000, 0.000, 1.000]

PE[1] = [sin(1·1.0), cos(1·1.0), sin(1·0.1), cos(1·0.1), ...]
       = [0.841,      0.540,      0.0998,      0.995,      0.00999, 0.99995, ...]

PE[2] = [sin(2·1.0), cos(2·1.0), sin(2·0.1), cos(2·0.1), ...]
       = [0.909,     -0.416,      0.198,      0.980,      0.01999, 0.99980, ...]

PE[3] = [sin(3·1.0), cos(3·1.0), sin(3·0.1), cos(3·0.1), ...]
       = [0.141,     -0.990,      0.296,      0.955,      0.02999, 0.99955, ...]
```

Notice: dimension 0-1 (high frequency) changes dramatically between positions. Dimension 6-7 (low frequency) barely changes. Together they form a *unique fingerprint* for each position.

The final input matrix $\tilde{X} = X + PE$:

```
X[t] = embed_vector[t] + PE[t]

For t=0: [-0.90, 0.10, 0.20, -0.30, ...]
         + [ 0.00, 1.00, 0.00,  1.00, ...]
         = [-0.90, 1.10, 0.20,  0.70, ...]
```

Figure [Figure 5.1](#) shows the sinusoidal positional encoding values as a heatmap over positions and dimensions.




Figure 5.1: Sinusoidal positional encoding heatmap

5.5 The Code: Positional Encoding in Python

The matrix data abstraction from Chapter 2 carries over unchanged.

```
# tag::make_matrix[]
def make_matrix(rows: int, cols: int, fill: float = 0.0) -> Matrix:
    return [[fill for _ in range(cols)] for _ in range(rows)]

def shape(matrix: Matrix) -> tuple[int, int]:
    return len(matrix), len(matrix[0]) if matrix else 0
# end::make_matrix[]

# tag::random_matrix[]
def random_matrix(
    rows: int,
    cols: int,
    rng: random.Random,
    scale: float | None = None,
) -> Matrix:
```



```

    scale = math.sqrt(2.0 / (rows + cols)) if scale is None else scale
    return [[rng.uniform(-scale, scale) for _ in range(cols)] for _ in range(rows)]
# end::random_matrix[]

```

`sinusoidal_position` encodes one position as a vector of alternating sine and cosine values. Dimension pair k oscillates at frequency $\omega_k = 1/10000^{2k/d}$, so each position gets a unique fingerprint.

```

def sinusoidal_position(position: int, d_model: int) -> Vector:
    row = []
    for i in range(d_model):
        k = i // 2
        omega = 1.0 / (10000.0 ** (2.0 * k / d_model))
        row.append(math.sin(position * omega) if i % 2 == 0 else math.cos(position * omega))
    return row

```

`sinusoidal_encoding` builds the full $[T \times d]$ matrix by stacking one row per position.

```

def sinusoidal_encoding(max_length: int, d_model: int) -> Matrix:
    return [sinusoidal_position(position, d_model) for position in range(max_length)]

```

`add_positional_encoding` sums the token embedding matrix X with the PE matrix element-wise, giving $\tilde{X} = X + PE$.

```

def add_positional_encoding(tokens: Matrix, positions: Matrix) -> Matrix:
    return matrix_add(tokens, positions[: len(tokens)])

```

Demo. Run with `python3 src/python/chapter_demos.py`:

```

def chapter_05() -> dict[str, object]:
    encoding = sinusoidal_encoding(4, 6)
    return {
        "shape": (len(encoding), len(encoding[0])),
        "first": encoding[0],
        "second": encoding[1],
    }

```

Figure Figure 5.2 shows the positional encoding matrix whose rows are added to token embeddings.

Figure Figure 5.3 shows how sinusoidal waves at different frequencies encode positions across dimensions.




Figure 5.2: Positional encoding matrix




Figure 5.3: Sinusoidal positional frequencies

5.6 Learned vs Sinusoidal Positional Embeddings

GPT-2 and GPT-3 use *learned* positional embeddings — a second embedding matrix $P \in \mathbb{R}^{T_{max} \times d}$, trained alongside the token embedding matrix.

Pros:

- Can adapt to the specific patterns in the training data.

Cons:

- Cannot generalize to sequences longer than T_{max} (the training context length).

Sinusoidal encoding is deterministic and extrapolates naturally. Modern architectures (LLaMA, Mistral) use a more sophisticated variant called *Rotary Positional Encoding (RoPE)* which applies the positional information inside the attention computation rather than at the input stage.

5.7 Key Takeaways

- Without position information, the model is *permutation-invariant* — it cannot distinguish "dog bites man" from "man bites dog".
- Sinusoidal PE encodes position using sine/cosine at exponentially-spaced frequencies.
- The encoding is *added* to the token embedding: $\tilde{X} = X + PE$.
- Nearby positions have high cosine similarity; distant positions have lower similarity.
- Modern models (GPT-2: learned; LLaMA: RoPE) vary the method but the purpose is the same.

Note

What's next? We now have $\tilde{X} \in \mathbb{R}^{T \times d}$ — a matrix that knows both what each token is and where it sits. Now the interesting part begins: *attention*, the mechanism that lets tokens share information with each other. Chapter 4.

Part III

Transformer Core

This part builds the transformer block from its major components. The chapters move from one attention head to a full block that can be stacked into a GPT model.

- In Chapter 6, each token learns which other tokens to read from.
- In Chapter 7, several attention patterns run in parallel.
- In Chapter 8, each token representation is transformed by a position-wise neural network.
- In Chapter 9, attention, feed-forward layers, residual connections, and normalization become one reusable block.

6 Attention — Tokens Talking to Each Other

Attention is the engine of the transformer. It is the mechanism that lets each token look at all (or some) of the other tokens and decide which ones to borrow information from. Without attention, each token would be processed independently, oblivious to context. With attention, the word “bank” can tell whether the surrounding tokens are about rivers or finance.

This chapter builds scaled dot-product attention from first principles.

6.1 The Idea

Imagine you are at a library. You arrive with a *query*: “I’m looking for books about Renaissance painting.” Every book in the library has a *key* on its spine — a short descriptor of its contents. You compare your query against every key, giving each book a *score*. Then you take a weighted blend of the books’ *values* (actual content) — spending more time on high-scoring books.

In the attention mechanism:

- *Query (Q)*: “What am I looking for?” — produced from the current token.
- *Key (K)*: “What do I contain?” — produced from every token in the sequence.
- *Value (V)*: “What information can I give?” — produced from every token in the sequence.

The output for each token is a *weighted sum* of all value vectors, where the weights come from comparing that token’s query against every token’s key.

Note

Math Minute — Softmax

`softmax(z)` turns a vector of raw scores into a probability distribution. For vector $z = [z_1, z_2, \dots, z_n]$:

6.2 [source]

6.3 $\text{softmax}(\mathbf{z}) = \exp(\mathbf{z}) / \sum \exp(\mathbf{z})$

Output values are all positive and sum to 1. Large scores get amplified; small scores get suppressed. The largest score “wins” most of the weight.

6.4 The Math

6.4.1 Step 1: Project to Q, K, V

Given input $X \in \mathbb{R}^{T \times d}$, we learn three weight matrices:

$$W_q \in \mathbb{R}^{d \times d_k}, \quad W_k \in \mathbb{R}^{d \times d_k}, \quad W_v \in \mathbb{R}^{d \times d_v}$$

And compute:

$$Q = XW_q \in \mathbb{R}^{T \times d_k} \quad K = XW_k \in \mathbb{R}^{T \times d_k} \quad V = XW_v \in \mathbb{R}^{T \times d_v}$$

In standard single-head attention, $d_k = d_v = d$.

6.4.2 Step 2: Compute Attention Scores

$$S = QK^\top / \sqrt{d_k} \in \mathbb{R}^{T \times T}$$

$S[i, j]$ measures how relevant token j is to token i . The division by $\sqrt{d_k}$ prevents the dot products from growing too large.

i Note

Math Minute — Matrix Multiplication

For matrices $A \in \mathbb{R}^{m \times n}$ and $B \in \mathbb{R}^{n \times p}$: $(AB)[i, j] = \sum_k A[i, k] \cdot B[k, j]$. The result has shape $[m \times p]$. Each entry is a dot product of a row of **A** and a column of **B**.

6.4.3 Step 3: Apply Causal Mask (for autoregressive models)

GPT is *autoregressive*: when predicting token t , it must not look at tokens $t+1$, $t+2$, We enforce this by masking the upper triangle:

$$M[i, j] = \begin{cases} 0 & \text{if } j \leq i \\ -\infty & \text{if } j > i \end{cases}$$

$$S_{\text{masked}} = S + M$$

Adding $-\infty$ before softmax effectively zeroes out those attention weights.

6.4.4 Step 4: Softmax $\rightarrow$ Attention Weights

$$A = \text{softmax}(S_{\text{masked}}) \in \mathbb{R}^{T \times T}$$

$A[i, j]$ is now a probability: how much token i attends to token j .

6.4.5 Step 5: Weighted Sum of Values

$$\text{Output} = AV \in \mathbb{R}^{T \times d_v}$$

$\text{Output}[i]$ is a weighted blend of all value vectors, guided by how much token i attends to each position.

Full formula:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}} + M\right) \cdot V$$

6.5 The Matrix: Worked Example

Let $T = 3$ tokens, $d = 4$, $d_k = d_v = 4$. Input (position-encoded):

```
X = [[ 1.0,  0.0,  1.0,  0.0],  ← token 0: "the"
      [ 0.0,  1.0,  0.0,  1.0],  ← token 1: "cat"
      [ 1.0,  1.0,  0.0,  0.0]]  ← token 2: "sat"
```

Weight matrices (simplified, identity-like): $W_q = W_k = W_v = I_4$ (4×4 identity, so $Q=K=V=X$).

Step 1 — Q, K, V (same as X here):

```
Q = K = V = X
```

Step 2 — Scores $S = QK^\top / \sqrt{4}$:

```
QK [0,0] = [1,0,1,0] · [1,0,1,0] = 2
QK [0,1] = [1,0,1,0] · [0,1,0,1] = 0
QK [0,2] = [1,0,1,0] · [1,1,0,0] = 1
```

```
S = QK / 2 =
  [[1.0,  0.0,  0.5],
   [0.0,  1.0,  0.5],
   [0.5,  0.5,  1.0]]
```

Step 3 — Causal mask:

```
S_masked =
  [[1.0,  -∞,  -∞],  ← token 0 can only attend to itself
   [0.0,  1.0,  -∞],  ← token 1 can attend to 0 and 1
   [0.5,  0.5,  1.0]] ← token 2 can attend to all
```

Step 4 — Softmax:

```
A[0] = softmax([1.0, -∞, -∞]) = [1.000, 0.000, 0.000]
A[1] = softmax([0.0, 1.0, -∞]) = [0.269, 0.731, 0.000]
A[2] = softmax([0.5, 0.5, 1.0]) = [0.211, 0.211, 0.578]
```

Step 5 — Output = AV :

```

Output[0] = 1.000 · V[0] = [1.0, 0.0, 1.0, 0.0]
Output[1] = 0.269 · [1,0,1,0] + 0.731 · [0,1,0,1] = [0.269, 0.731, 0.269, 0.731]
Output[2] = 0.211 · [1,0,1,0] + 0.211 · [0,1,0,1] + 0.578 · [1,1,0,0]
          = [0.789, 0.789, 0.211, 0.211]

```

Token 2 (“sat”) is blending information from all three tokens, with most weight on itself.

Figure Figure 6.1 shows Q, K, and V as three separate linear projections of the same input matrix X.




Figure 6.1: Query, key, and value projections

Figure Figure 6.2 shows causal attention weights, where rows are queries, columns are keys, and future tokens are masked.

6.6 The Code: Scaled Dot-Product Attention in Python

Attention is a composition of named operations. The matrix abstraction from Chapter 2 carries forward; we add multiply, transpose, scale, and add.

```

# tag::transpose[]
def transpose(matrix: Matrix) -> Matrix:
    return [list(col) for col in zip(*matrix)]
# end::transpose[]

```




Figure 6.2: Causal attention weight heatmap

```

def matrix_add(left: Matrix, right: Matrix) -> Matrix:
    return [[a + b for a, b in zip(row_a, row_b)] for row_a, row_b in zip(left, right)]

def matrix_scale(matrix: Matrix, scalar: float) -> Matrix:
    return [[scalar * value for value in row] for row in matrix]

# tag::matrix_multiply[]
def matrix_multiply(left: Matrix, right: Matrix) -> Matrix:
    right_t = transpose(right)
    return [[dot(row, col) for col in right_t] for row in left]
# end::matrix_multiply[]

def hstack(matrices: Sequence[Matrix]) -> Matrix:
    return [sum((matrix[row] for matrix in matrices), []) for row in range(len(matrices[0]))]

```

`softmax` converts a vector of raw scores into a probability distribution. We subtract the maximum value before exponentiating to prevent overflow.

```

# tag::softmax_vector[]
def softmax(logits: Sequence[float]) -> Vector:
    max_logit = max(logits)
    exp_values = [math.exp(value - max_logit) for value in logits]
    total = sum(exp_values)
    return [value / total for value in exp_values]
# end::softmax_vector[]

# tag::softmax_rows[]
def softmax_rows(matrix: Matrix) -> Matrix:
    return [softmax(row) for row in matrix]
# end::softmax_rows[]

```

`causal_mask` fills the upper triangle with -10^9 so future tokens cannot be attended to.

```

def causal_mask(size: int) -> Matrix:
    return [[0.0 if j <= i else -1.0e9 for j in range(size)] for i in range(size)]

```

scaled_dot_product_attention and self_attention are direct translations of the attention formula.

```
def scaled_dot_product_attention(query: Matrix, key: Matrix, value: Matrix) -> tuple[Matrix,
    d_key = shape(query)[1]
    scores = matrix_scale(matrix_multiply(query, transpose(key)), 1.0 / math.sqrt(d_key))
    masked_scores = matrix_add(scores, causal_mask(len(scores)))
    weights = softmax_rows(masked_scores)
    return matrix_multiply(weights, value), weights

def self_attention(x: Matrix, wq: Matrix, wk: Matrix, wv: Matrix) -> tuple[Matrix, Matrix]:
    return scaled_dot_product_attention(
        matrix_multiply(x, wq),
        matrix_multiply(x, wk),
        matrix_multiply(x, wv),
    )
```

Run with `python3 src/python/chapter_demos.py`.

6.7 Key Takeaways

- Attention asks: for each token, *which other tokens are most relevant?*
- Q, K, V are *learned linear projections* of the input.
- The attention score $S[i, j] = Q_i \cdot K_j / \sqrt{d_k}$ measures relevance.
- Causal masking prevents tokens from seeing future positions (critical for text generation).
- The output is a *soft, weighted average* of value vectors.
- Full formula: $\text{Attention}(Q, K, V) = \text{softmax}(QK^\top / \sqrt{d_k} + \text{mask})V$.

Note

What's next? One attention head can only track one type of relationship at a time. What if we want to simultaneously track syntactic dependencies, coreference, and semantic similarity? We run *multiple attention heads in parallel* — Chapter 5.

7 Multi-Head Attention — Many Conversations at Once

A single attention head is expressive, but limited. It produces one set of attention weights — one pattern of “who attends to whom.” In natural language, multiple distinct relationships coexist in the same sentence. Consider:

```
"The animal didn't cross the street because it was too tired."
```

- One pattern might link “it” → “animal” (coreference)
- Another might link “cross” → “street” (verb-object)
- Another might track the causal structure “because”

Multi-head attention runs H independent attention heads in parallel, each free to specialize on a different relationship type. Their outputs are concatenated and projected back to the model dimension.

7.1 The Idea

A single attention operation produces one pattern of “who attends to whom.” But a sentence carries many different kinds of relationships at the same time.

In “The animal didn’t cross the street because it was too tired”:

- *it* refers back to *animal* — that is a *coreference* relationship.
- *cross* links to *street* — that is a *verb-object* relationship.
- *because* connects a cause to an effect — that is a *logical* relationship.

One attention head can only focus on one of these at a time. Multi-head attention runs several independent attention operations *in parallel* — each one free to specialize on a different pattern. Each head sees the same input but learns to ask a different question of it.

At the end, the results from all heads are stitched back together and projected into a single vector, the same size as before. The model learns entirely from data which head should track grammar, which should track meaning, which should track proximity — no one programs this in explicitly.

The result is a richer representation than any single head could produce: each token's final vector has been informed by multiple different lenses simultaneously.

7.2 The Math

Step 1 — Compute each head independently:

For $h = 1, \dots, H$:

$$\begin{aligned} Q^h &= XW_q^h \in \mathbb{R}^{T \times d_k} \\ K^h &= XW_k^h \in \mathbb{R}^{T \times d_k} \\ V^h &= XW_v^h \in \mathbb{R}^{T \times d_v} \\ \text{head}^h &= \text{softmax}\left(\frac{Q^h K^{h\top}}{\sqrt{d_k}} + M\right) V^h \end{aligned}$$

Step 2 — Concatenate:

$$\text{MultiHead} = \text{concat}(\text{head}^1, \text{head}^2, \dots, \text{head}^H) \in \mathbb{R}^{T \times (H \cdot d_v)}$$

Since $d_v = d/H$, the concatenated output has shape $[T \times d]$ — the same as the input.

Step 3 — Output projection:

$W_o \in \mathbb{R}^{d \times d}$ is a learned output projection matrix.

$$\text{Output} = \text{MultiHead} \cdot W_o \in \mathbb{R}^{T \times d}$$

The final output projection mixes information across heads.

Full formula:

$$\text{MHA}(X) = [\text{head}^1 \parallel \text{head}^2 \parallel \dots \parallel \text{head}^H] W_o$$

where $\text{head}^h = \text{Attention}(XW_q^h, XW_k^h, XW_v^h)$.

7.3 The Matrix: Worked Example

Let $T = 3$, $d = 4$, $H = 2$ heads, so $d_k = d_v = 2$.

Input:

$$X = \begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 0 & 0 \end{bmatrix} \quad (3 \times 4)$$

Head 1 uses the first 2 dimensions primarily.

$$Wq^1 = Wk^1 = Wv^1 = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} \quad (4 \times 2)$$

$$Q^1 = X Wq^1 = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 1 \end{bmatrix} \quad (3 \times 2)$$

Scores: $S^1 = Q^1 K^{1\top} / \sqrt{2}$:

$$\begin{aligned} Q^1 K^1 &= \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 1 & 2 \end{bmatrix} \\ S^1 &= \begin{bmatrix} 0.71 & 0.00 & 0.71 \\ 0.00 & 0.71 & 0.71 \\ 0.71 & 0.71 & 1.41 \end{bmatrix} \end{aligned}$$

After causal mask and softmax:

$$A^1 = \begin{bmatrix} 1.000 & 0.000 & 0.000 \\ 0.414 & 0.586 & 0.000 \\ 0.221 & 0.221 & 0.558 \end{bmatrix}$$

$$\text{head}^1 = A^1 V^1 = \begin{bmatrix} 1.000 & 0.000 \\ 0.586 & 0.414 \\ 0.779 & 0.779 \end{bmatrix} \quad (3 \times 2)$$

Head 2 focuses on last 2 dimensions, producing similarly shaped output. After concatenating both heads and applying the output projection $W_o \in \mathbb{R}^{4 \times 4}$:

$$\begin{aligned} \text{MultiHead} &= [\text{head}^1 \parallel \text{head}^2] = \\ &= \begin{bmatrix} 1.000 & 0.000 & 1.000 & 0.000 \\ 0.586 & 0.414 & 0.500 & 0.500 \\ 0.779 & 0.779 & 0.421 & 0.211 \end{bmatrix} \quad (3 \times 4) \end{aligned}$$

Figure Figure 7.1 shows parallel attention heads whose outputs are concatenated and projected back to model width.




Figure 7.1: Multi-head attention

7.4 The Code: Multi-Head Attention in Python

Each attention head is a triple of weight matrices (W_q, W_k, W_v) , each of shape $[d \times d_k]$. `MultiHeadAttention` groups the heads with the output projection W_o ; `make_multi_head_attention` allocates them.

```

@dataclass
class AttentionHead:
    wq: Matrix
    wk: Matrix
    wv: Matrix

@dataclass
class MultiHeadAttention:
    heads: list[AttentionHead]
    wo: Matrix

def make_multi_head_attention(d_model: int, num_heads: int, rng: random.Random) -> MultiHeadAttention:
    if d_model % num_heads != 0:
        raise ValueError("d_model must be divisible by num_heads")
    d_key = d_model // num_heads
    heads = [

```

```

    AttentionHead(
        random_matrix(d_model, d_key, rng),
        random_matrix(d_model, d_key, rng),
        random_matrix(d_model, d_key, rng),
    )
    for _ in range(num_heads)
]
return MultiHeadAttention(heads=heads, wo=random_matrix(d_model, d_model, rng))

```

`multi_head_attention` runs each head's SDPA independently, concatenates the per-head outputs column-wise, and applies the output projection.

```

def multi_head_attention(x: Matrix, params: MultiHeadAttention) -> tuple[Matrix, list[Matrix]]:
    results = [
        self_attention(x, head.wq, head.wk, head.wv)
        for head in params.heads
    ]
    concatenated = hstack([output for output, _weights in results])
    return matrix_multiply(concatenated, params.wo), [weights for _output, weights in results]

```

Demo. Run with `python3 src/python/chapter_demos.py`:

```

def chapter_07(seed: int = 7) -> dict[str, object]:
    rng = random.Random(seed)
    x = random_matrix(4, 8, rng)
    params = make_multi_head_attention(8, 2, rng)
    output, weights = multi_head_attention(x, params)
    return {
        "output_shape": (len(output), len(output[0])),
        "num_heads": len(weights),
    }

```

7.5 Why Multi-Head Attention Works

Each head learns to specialize. Research has identified heads that:

- Track *syntactic structure* (subject-verb agreement)
- Resolve *coreference* (“it” → “the cat”)
- Handle *positional offsets* (“look 2 tokens back”)
- Track *rare-word semantics*

These specializations emerge from training, not from explicit design.

i Note

Math Minute — Expressivity

H heads of dimension d/H can represent attention patterns that a single head of dimension d cannot easily learn. This is analogous to having H different “lenses” looking at the same sequence — each lens focuses on different features. The output projection W_o then combines the views.

7.6 Key Takeaways

- Multi-head attention runs H attention heads *in parallel*, each in a lower-dimensional subspace $d_k = d/H$.
- Each head has independent W_q, W_k, W_v matrices — each learns a different “question to ask.”
- Outputs are *concatenated* (not averaged), then projected back to d with W_o .
- The total parameter count is $3Hd \cdot d_k + d^2 = 3d^2 + d^2$ — same as one large head.
- Different heads specialize in different linguistic relationships.

i Note

What’s next? After attention mixes information across tokens, each token’s vector goes through a small *feed-forward network* — a two-layer MLP applied identically to every position. This is where most of the model’s stored “knowledge” lives. Chapter 6.

8 Feed-Forward Network — The Model’s Memory

After multi-head attention has blended information across positions, each token’s vector passes through a *feed-forward network (FFN)*. This is a simple two-layer MLP — the same network, applied independently to every position.

While attention does the *routing* (mixing information across tokens), the FFN does the *processing* (transforming each token’s representation). Experiments by Geva et al. (2021) showed that the FFN layers act as “key-value memories” — the first layer’s neurons trigger on specific input patterns, and the second layer retrieves associated information.

8.1 The Idea

Attention is about *communication* — each token collects information from other tokens. The feed-forward network is about *computation* — each token processes what it just collected, entirely on its own.

After attention, every token’s vector has been updated with context from the surrounding words. The FFN now takes each updated vector and runs it through a private transformation: expand it into a much larger space, filter it through a non-linearity that decides which features are “active,” then compress it back to the original size. No information crosses between tokens here — every position is handled independently with the same set of weights.

Why expand and then compress? The expansion gives the model room to express a large number of potential features at once. The non-linearity then selects which combinations matter and switches the rest off. The compression packages the result back into the model’s standard vector size.

This expand-filter-compress pipeline is where most of the model’s factual knowledge lives. Research has shown that individual neurons in the FFN’s expanded space can be associated with specific concepts (“Paris,” “past tense,” “chemical element”).

The same FFN weights are applied to every token position in the sequence: position 1 and position 100 go through identical transformations. Only attention sees position.

8.2 The Math

i Note

Math Minute — ReLU and GELU

A *non-linearity* is a function that makes the network capable of learning complex patterns beyond linear mappings.

$\text{ReLU}(\mathbf{x}) = \max(0, \mathbf{x})$ — simple: negative values $\rightarrow 0$, positive values unchanged.

$\text{GELU}(x) = x \cdot \Phi(x)$ where Φ is the standard normal CDF. GELU is smooth and probabilistic. GPT uses GELU. LLaMA and most modern models use **SwiGLU**, a gated variant.

$\text{Sigmoid}(\mathbf{x}) = 1/(1+e^{-\mathbf{x}})$ squashes any real number to $(0,1)$. Used in gates.

For a sequence input $X \in \mathbb{R}^{T \times d}$, the FFN applies identically to each row:

$$\text{FFN}(X) = \text{GELU}(XW_1 + \mathbf{1} \cdot b_1^\top)W_2 + \mathbf{1} \cdot b_2^\top$$

Breaking it down:

$$H = XW_1 + \mathbf{1}b_1^\top \in \mathbb{R}^{T \times d_{\text{ff}}} \quad (\text{expand})$$

$$H' = \text{GELU}(H) \in \mathbb{R}^{T \times d_{\text{ff}}} \quad (\text{activate})$$

$$Y = H'W_2 + \mathbf{1}b_2^\top \in \mathbb{R}^{T \times d} \quad (\text{contract})$$

Parameter count:

- W_1 : $d \times d_{\text{ff}} = d \times 4d = 4d^2$
- W_2 : $d_{\text{ff}} \times d = 4d^2$
- Total: $\approx 8d^2$ — *twice* the parameter count of all attention weight matrices combined!

8.3 GELU Deep Dive

The exact GELU formula is:

$$\text{GELU}(x) = x \cdot \frac{1}{2} \cdot \left[1 + \text{erf}\left(\frac{x}{\sqrt{2}}\right) \right]$$

A fast approximation (used in practice):

$$\text{GELU}(x) \approx 0.5 \cdot x \cdot (1 + \tanh(\sqrt{2/\pi} \cdot (x + 0.044715 \cdot x^3)))$$

Key behaviors:

- $\text{GELU}(0) = 0$
- For large positive x : $\text{GELU}(x) \approx x$ (passes through)
- For large negative x : $\text{GELU}(x) \approx 0$ (suppressed)
- Smooth everywhere — gradient flows through during training

8.4 The Matrix: Worked Example

Let $T = 2$, $d = 4$, $d_{\text{ff}} = 8$.

Input (after attention):

```
X = [[ 0.5,  1.2, -0.3,  0.8],  
      [-0.1,  0.7,  0.9, -0.4]]    (2×4)
```

Let's trace just row 0 of X: $x = [0.5, 1.2, -0.3, 0.8]$

Step 1 — Expand: $h = xW_1$

```
h = [1.1, -0.4, 0.8, 2.1, -0.3, 0.6, 1.5, -0.8]
```

Step 2 — GELU:

```
GELU(1.1)    0.95  
GELU(-0.4)   -0.12  
GELU(0.8)    0.64  
GELU(2.1)    2.06  
GELU(-0.3)   -0.11  
GELU(0.6)    0.44  
GELU(1.5)    1.40  
GELU(-0.8)   -0.17
```

```
h' = [0.95, -0.12, 0.64, 2.06, -0.11, 0.44, 1.40, -0.17]
```

Step 3 — Contract: $y = h'W_2$

```
y  [0.7, 1.1, -0.2, 0.9]
```




Figure 8.1: Feed-forward network

The output y is a transformed version of the input vector — same dimensionality, different values. The transformation was learned to improve next-token prediction.

Figure Figure 8.1 shows the feed-forward network expanding the hidden dimension, applying GELU, and contracting back to model width.

Figure Figure 8.2 shows the FFN interpretation as associative memory, with keys and values stored in its weight matrices.

8.5 The Code: FFN in Python

GELU is the activation function used in GPT-2. For large positive x , GELU approaches x — the value passes through. For large negative x , it approaches 0 — the value is suppressed. `gelu_matrix` applies the scalar GELU to every element of a matrix.

```
def sigmoid(x: float) -> float:
    return 1.0 / (1.0 + math.exp(-x))

def relu(x: float) -> float:
    return max(0.0, x)
```




Figure 8.2: Feed-forward network as associative memory

```
def gelu(x: float) -> float:
    scale = math.sqrt(2.0 / math.pi)
    return 0.5 * x * (1.0 + math.tanh(scale * (x + 0.044715 * x**3)))

def gelu_matrix(matrix: Matrix) -> Matrix:
    return [[gelu(value) for value in row] for row in matrix]
```

`add_bias` broadcasts a bias vector b across all rows of a matrix M .

```
def add_bias(matrix: Matrix, bias: Vector) -> Matrix:
    return [[value + bias[j] for j, value in enumerate(row)] for row in matrix]
```

`FeedForward` holds the four parameters; `make_feed_forward` allocates them — W_1 expands the model dimension d to $4d$; W_2 contracts it back.

```
@dataclass
class FeedForward:
    w1: Matrix
    b1: Vector
```



```

w2: Matrix
b2: Vector

def make_feed_forward(d_model: int, rng: random.Random) -> FeedForward:
    d_hidden = 4 * d_model
    return FeedForward(
        w1=random_matrix(d_model, d_hidden, rng),
        b1=[0.0] * d_hidden,
        w2=random_matrix(d_hidden, d_model, rng),
        b2=[0.0] * d_model,
    )

```

`feed_forward` composes the two linear layers with GELU: expand $\rightarrow$ activate $\rightarrow$ contract.

```

def feed_forward(x: Matrix, params: FeedForward) -> Matrix:
    hidden = gelu_matrix(add_bias(matrix_multiply(x, params.w1), params.b1))
    return add_bias(matrix_multiply(hidden, params.w2), params.b2)

```

Demo. Run with `python3 src/python/chapter_demos.py`:

```

def chapter_08(seed: int = 8) -> dict[str, object]:
    rng = random.Random(seed)
    x = random_matrix(3, 8, rng)
    params = make_feed_forward(8, rng)
    output = feed_forward(x, params)
    return {"output_shape": (len(output), len(output[0]))}

```

8.6 The FFN as a Key-Value Memory

A striking insight from Geva et al. (2021): the FFN layers act like *associative memories*.

The first layer W_1 stores “keys” — each column of W_1 is a pattern to match against the input. GELU fires when the input matches. The second layer W_2 stores “values” — the corresponding information to retrieve.

Concretely:

$$h = \text{GELU}(xW_1)$$

where each component $h_k \approx x \cdot W_1[:, k]$ — a dot-product “match score” for pattern k . High h_k means “input matches pattern k .”

$$y = hW_2$$

Pull out value $W_2[k, :]$ weighted by h_k .

This means the model can learn: “if the input contains patterns related to France, activate neuron 47, which retrieves ‘Paris’ associations.”

8.7 Key Takeaways

- The FFN is a two-layer MLP applied *identically to every position*: $FFN(x) = \text{GELU}(xW_1 + b_1)W_2 + b_2$.
- The inner dimension $d_{\text{ff}} = 4d$ means the FFN has *more parameters than the attention layers*.
- GELU is a smooth activation function that “softly gates” negative values.
- The FFN acts as an *associative memory*: first layer matches patterns, second layer retrieves associated information.
- No cross-position communication happens here — that’s attention’s job.

i Note

What’s next? We now have all the pieces: attention that mixes information, and an FFN that processes each token’s vector. How are they wired together? With *residual connections* and *layer normalization* to form the *transformer block* — Chapter 7.

9 The Transformer Block — Putting It Together

A transformer block combines multi-head attention and the feed-forward network with two crucial wiring techniques: *residual connections* and *layer normalization*. Without these, deep stacks of transformer blocks would fail to train — gradients would vanish or explode.

A GPT model is simply a stack of N identical transformer blocks. GPT-2 small: 12 blocks. GPT-3 175B: 96 blocks.

9.1 The Idea

A GPT model is a stack of identical blocks — GPT-2 small has 12, GPT-3 has 96. Stacking many layers should make the model more powerful, but in practice it causes two problems: information gets distorted as it passes through many transformations, and very deep networks are notoriously hard to train.

A transformer block solves both problems with two simple ideas.

Residual connections (also called skip connections): instead of replacing a vector entirely, each sub-layer only adds its output to the original. The input is preserved and combined with the new information. Think of it as writing notes in the margin rather than rewriting the whole page. Each layer only needs to learn *what to change*, not what the whole answer should be. This keeps information flowing cleanly all the way through, even in very deep stacks.

Layer normalization: the numbers inside a vector can grow or shrink unpredictably as the model trains. If they drift too far, the training process becomes unstable. Layer normalization resets them to a consistent scale before each sub-layer processes them. It is a housekeeping step: it does not change which direction the vector points, just how large the numbers are.

Together, these two techniques make it possible to stack dozens of transformer blocks without losing control of training.

9.2 The Math

A single transformer block (Pre-LN variant, used by most modern GPTs):

$$\begin{aligned}x_1 &= x + \text{MHA}(\text{LayerNorm}(x)) && \text{(attention sub-layer)} \\x_2 &= x_1 + \text{FFN}(\text{LayerNorm}(x_1)) && \text{(FFN sub-layer)}\end{aligned}$$

This is the *Pre-Layer-Norm* architecture. The original “Attention Is All You Need” paper used Post-LN, but Pre-LN is more stable to train and is used in GPT-2 onwards.

9.2.1 Layer Normalization

For a vector $x \in \mathbb{R}^d$:

$$\text{LayerNorm}(x) = \gamma \odot \frac{x - \mu}{\sigma} + \beta$$

where $\mu = \frac{1}{d} \sum_i x_i$, $\sigma = \sqrt{\frac{1}{d} \sum_i (x_i - \mu)^2}$, $\gamma, \beta \in \mathbb{R}^d$ are learned scale and shift, and $\odot$ is element-wise multiplication.

Note

Math Minute — Variance and Standard Deviation

The variance of a set of numbers $\{x_1, \dots, x_n\}$ measures how spread out they are: $Var = (1/n) \sum_i (x_i - \mu)^2$ where μ is the mean. The standard deviation $\sigma = \sqrt{Var}$ is in the same units as the original values. Dividing by σ makes the spread equal to 1 — “standardizing.”

9.3 The Residual Stream

A powerful way to understand the GPT architecture is through the lens of the *residual stream* (Elhage et al., Anthropic 2021):

The input embedding is injected into a “stream” — a vector of dimension d per token. Each transformer block *reads from* this stream (via attention and FFN) and *adds back to* it (via residual connections). The stream carries information across all blocks.

```

stream = token_embeddings + positional_encodings    [T × d]
stream1 = stream + MHA(LN(stream))
stream1 = stream1 + FFN(LN(stream1))
stream2 = stream1 + MHA(LN(stream1))
stream2 = stream2 + FFN(LN(stream2))

stream = final output

```

This view clarifies that attention heads and FFN neurons are all writing to the same shared workspace.

9.4 The Matrix: Worked Example

Let $T = 2, d = 4$.

Input embedding (position-encoded):

```
x = stream[0] = [1.0, -0.5, 0.8, -0.2]
```

LayerNorm step:

```

= (1.0 - 0.5 + 0.8 - 0.2) / 4 = 0.275
2 = [(1.0-0.275)2 + (-0.5-0.275)2 + (0.8-0.275)2 + (-0.2-0.275)2] / 4
    = 0.407
    = √0.407    0.638

x_norm = [1.136, -1.215, 0.823, -0.745]
(With =1, =0 for simplicity)

```

MHA produces (let's say): `mha_out` = [0.3, 0.7, -0.1, 0.5]

Residual connection:


```

x = x + mha_out = [1.0+0.3, -0.5+0.7, 0.8-0.1, -0.2+0.5]
  = [1.3, 0.2, 0.7, 0.3]

```

Second LayerNorm + FFN + residual (analogously) $\rightarrow x_2$.

The key insight: x_1 contains *both* the original information (from `x`) and the new information (from `mha_out`). Nothing is overwritten — the residual stream accumulates.



Transformer Block (stacked N times)

Figure 9.1: Pre-LN transformer block

Figure 9.1 shows the full Pre-LN transformer block with LayerNorm before each sub-layer and residual connections around both.

Figure 9.2 shows the residual stream carrying information forward while gradients flow backward through additions.




Figure 9.2: Residual stream

9.5 The Code: Transformer Block in Python

Layer normalization shifts a single vector to mean 0, variance 1, then applies learned scale and shift. `LayerNorm` holds the learned parameters γ and β ; `make_layer_norm` initializes them to identity.

```
@dataclass
class LayerNorm:
    gamma: Vector
    beta: Vector
    epsilon: float = 1.0e-5

def make_layer_norm(d_model: int) -> LayerNorm:
    return LayerNorm(gamma=[1.0] * d_model, beta=[0.0] * d_model)

def layer_norm_vector(vector: Vector, params: LayerNorm) -> Vector:
    avg = mean(vector)
    std = math.sqrt(variance(vector) + params.epsilon)
    return [
        params.gamma[i] * ((value - avg) / std) + params.beta[i]
```

```

    for i, value in enumerate(vector)
]

```

`layer_norm_matrix` applies `layer_norm_vector` independently to each row of a $[T \times d]$ matrix.

```

def layer_norm_matrix(matrix: Matrix, params: LayerNorm) -> Matrix:
    return [layer_norm_vector(row, params) for row in matrix]

```

`TransformerBlock` groups the four sub-layers; `make_transformer_block` allocates their parameters.

```

@dataclass
class TransformerBlock:
    attention: MultiHeadAttention
    feed_forward: FeedForward
    ln1: LayerNorm
    ln2: LayerNorm

def make_transformer_block(d_model: int, num_heads: int, rng: random.Random) -> TransformerBlock:
    return TransformerBlock(
        attention=make_multi_head_attention(d_model, num_heads, rng),
        feed_forward=make_feed_forward(d_model, rng),
        ln1=make_layer_norm(d_model),
        ln2=make_layer_norm(d_model),
    )

```

`transformer_block` is the Pre-LN residual block; `forward_stack` threads input through N blocks.

```

def transformer_block(x: Matrix, params: TransformerBlock) -> Matrix:
    attn_out, _weights = multi_head_attention(layer_norm_matrix(x, params.ln1), params.attention)
    x1 = matrix_add(x, attn_out)
    ffn_out = feed_forward(layer_norm_matrix(x1, params.ln2), params.feed_forward)
    return matrix_add(x1, ffn_out)

def forward_stack(x: Matrix, blocks: Iterable[TransformerBlock]) -> Matrix:
    for block in blocks:
        x = transformer_block(x, block)
    return x

```


Demo. Run with `python3 src/python/chapter_demos.py`:

```
def chapter_09(seed: int = 9) -> dict[str, object]:
    rng = random.Random(seed)
    x = random_matrix(3, 8, rng)
    block = make_transformer_block(8, 2, rng)
    normed = layer_norm_matrix(x, make_layer_norm(8))
    output = transformer_block(x, block)
    return {
        "normed_shape": (len(normed), len(normed[0])),
        "output_shape": (len(output), len(output[0])),
    }
```

9.6 Why N Blocks?

A single transformer block has limited capacity. Each block:

- Runs one round of attention (all heads) — *mixes information globally*
- Runs one FFN per position — *processes each token's updated representation*

But complex language tasks require many rounds of computation. By stacking N blocks, the model builds increasingly abstract representations:

- Early blocks: local patterns, syntax, common phrases
- Middle blocks: semantic roles, entity tracking
- Late blocks: task-specific reasoning, pragmatics

The residual stream carries information across all blocks. Later blocks can read from and write to everything earlier blocks computed.

9.7 Key Takeaways

- A transformer block = MHA $\rightarrow$ residual $\rightarrow$ LN $\rightarrow$ FFN $\rightarrow$ residual $\rightarrow$ LN (Pre-LN ordering).
- *Residual connections* let gradients flow directly and prevent vanishing gradients.
- *Layer normalization* keeps activations at mean 0, std 1 within each sub-layer.
- The *residual stream* is the shared workspace that all blocks read from and write to.
- GPT stacks N identical blocks; capacity scales with N and d .

i Note

What's next? After N blocks, we have a final matrix $X_{\text{final}} \in \mathbb{R}^{T \times d}$. Each row is a rich representation of the corresponding token in context. The last step: turn the final vector for position T into a probability distribution over the vocabulary — that's *vocabulary projection* in Chapter 8.

Part IV

Prediction and Learning

This part follows the model from final hidden states to training updates. It explains how GPT chooses the next token, measures mistakes, and changes its weights.

- In Chapter 10, the final vector becomes logits and next-token probabilities.
- In Chapter 11, predictions become a scalar training signal through cross-entropy and perplexity.
- In Chapter 12, gradients and weight updates connect the loss back to the model parameters.

10 Vocabulary Projection — From Vectors to Words

After N transformer blocks, we have a matrix $X_{\text{final}} \in \mathbb{R}^{T \times d}$. The final row, $X_{\text{final}}[T-1]$, is a rich contextualized representation of the last token. The last step is to project this vector into a probability distribution over the vocabulary: which token is most likely to come next?

This projection is called the *language model head*, or *unembedding layer*.

10.1 The Idea

After all the attention and feed-forward layers, we have a vector representing the last token — a list of hundreds of numbers that encodes everything the model knows about what has been said so far. But the user needs a *word*, not a list of numbers.

The final step converts that vector into a probability over the entire vocabulary: for each of the 50,000-odd words the model knows, what is the chance it is the right next word? The word with the highest probability (or a word sampled from the distribution) becomes the model’s output.

How does a vector become a probability distribution? The model compares the vector against every word in the vocabulary and assigns a score to each one. High score means “this word fits the context well.” Low score means it does not. Then those scores are squeezed into the range 0–1 and made to sum to 1, giving a proper probability.

This step is the mirror image of Chapter 2: embedding turned an integer (a word ID) into a vector; unembedding turns a vector back into a distribution over integers. Many implementations even reuse the same table for both directions — the same weights that encode “cat $\rightarrow$ vector” are transposed to decode “vector $\rightarrow$ how much does this look like cat?”

10.2 The Math

Given the final hidden state $h_t = X_{\text{final}}[T-1] \in \mathbb{R}^d$:

Step 1 — Final Layer Norm:

$$\tilde{h}_t = \text{LayerNorm}(h_t) \in \mathbb{R}^d$$

Step 2 — Vocabulary Projection:

$$\text{logits} = \tilde{h}_t W_u \in \mathbb{R}^{|V|}$$

With weight tying: $W_u = E^\top$, so:

$$\text{logits}[i] = \tilde{h}_t \cdot E_i$$

The logit for token i is the dot product of the model’s “prediction vector” with token i ’s embedding. Tokens whose embeddings align with the prediction vector get high logits.

Step 3 — Softmax $\rightarrow$ Probabilities:

$$P(\text{next} = i \mid \text{context}) = \text{softmax}(\text{logits})[i] = \frac{\exp(\text{logits}[i])}{\sum_j \exp(\text{logits}[j])}$$

i Note

Math Minute — Temperature

We can control the “sharpness” of the distribution with a *temperature parameter* T_{temp} :

$$P(i) = \text{softmax}(\text{logits}/T_{\text{temp}})[i]$$

- $T_{\text{temp}} \rightarrow 0$: argmax — always pick the highest-probability token (greedy)
- $T_{\text{temp}} = 1$: standard softmax
- $T_{\text{temp}} > 1$: flattens the distribution — more randomness/creativity

Temperature is not a model parameter — it’s a sampling hyperparameter set at inference time.

10.3 The Generation Loop

Training and inference use the same forward pass differently.

Training: Given a sequence $[t_1, t_2, \dots, t_n]$, the model predicts all next tokens simultaneously (thanks to causal masking): $P(t_2|t_1), P(t_3|t_1, t_2), \dots$. The loss is cross-entropy averaged over all positions.

Inference (generation):

1. Start with prompt tokens $[t_1, \dots, t_n]$
2. Forward pass $\rightarrow$ logits for position n
3. Sample or argmax $\rightarrow t_{n+1}$
4. Append t_{n+1} to the sequence
5. Forward pass $\rightarrow$ logits for position $n + 1$
6. Repeat until stop token or max length

This is called *autoregressive generation*: each new token is fed back in as input to generate the next.

10.4 The Matrix: Worked Example

Let $|V| = 5$, $d = 4$. Final hidden state:

```
h = [0.3, -0.1, 0.8, 0.2]
```

Unembedding matrix $W_u = E^\top$ (columns = token embeddings):

```
logits[0] = h · E[0] = (0.3)(0.10) + (-0.1)(-0.20) + (0.8)(0.30) + (0.2)(-0.40)
           = 0.030 + 0.020 + 0.240 - 0.080 = 0.210

logits[1] = h · E[1] = 0.150 - 0.060 - 0.560 + 0.160 = -0.310
logits[2] = h · E[2] = -0.270 - 0.010 + 0.160 - 0.060 = -0.180
logits[3] = h · E[3] = 0.120 + 0.050 + 0.480 - 0.140 = 0.510
logits[4] = h · E[4] = -0.030 - 0.080 - 0.320 + 0.100 = -0.330

logits = [0.210, -0.310, -0.180, 0.510, -0.330]
```

Softmax:

```
exp(logits) = [1.233, 0.733, 0.835, 1.665, 0.719]
sum          = 5.185

P = [0.238, 0.141, 0.161, 0.321, 0.139]
```

Token 3 has the highest probability (32.1%) — that's the model's best guess for next token.

Figure Figure 10.1 shows the vocabulary projection converting a hidden state into logits over the vocabulary.

Figure Figure 10.2 shows autoregressive generation, where each predicted token is appended and fed back into the model.




Figure 10.1: Vocabulary projection




Figure 10.2: Autoregressive generation loop

10.5 The Code: Full microGPT Forward Pass in Python

Configuration is data. We represent the model's five hyperparameters as a plain list and define selectors once.

```
@dataclass(frozen=True)
class GPTConfig:
    vocab_size: int
    d_model: int
    num_heads: int
    num_layers: int
    max_seq_len: int
```

`make_gpt_params` allocates all learnable weights. `wte` is the token embedding matrix $[|V| \times d]$; `wpe` is the position embedding matrix. Weight tying means the unembedding step will reuse `wte` transposed — no separate parameter is stored.

```
@dataclass
class GPTParams:
    wte: Matrix
    wpe: Matrix
    blocks: list[TransformerBlock]
    ln_f: LayerNorm

def make_gpt_params(config: GPTConfig, rng: random.Random) -> GPTParams:
    return GPTParams(
        wte=random_matrix(config.vocab_size, config.d_model, rng),
        wpe=random_matrix(config.max_seq_len, config.d_model, rng),
        blocks=[
            make_transformer_block(config.d_model, config.num_heads, rng)
            for _ in range(config.num_layers)
        ],
        ln_f=make_layer_norm(config.d_model),
    )
```

`gpt_forward` is the full pipeline: embed token IDs, add position embeddings, pass through N transformer blocks, apply the final layer norm to the last token's vector, then compute dot products with every row of `wte`.

```
def gpt_forward(token_ids: Sequence[int], params: GPTParams, config: GPTConfig) -> Vector:
    if len(token_ids) > config.max_seq_len:
        raise ValueError("token_ids exceeds max_seq_len")

    token_embeddings = embed_sequence(params.wte, token_ids)
    position_embeddings = embed_sequence(params.wpe, range(len(token_ids)))
    x = matrix_add(token_embeddings, position_embeddings)
    x = forward_stack(x, params.blocks)
    h = layer_norm_vector(x[-1], params.ln_f)
    return [dot(h, token_vector) for token_vector in params.wte]
```

`generate` is the autoregressive loop.

```
def temperature_scale(logits: Sequence[float], temperature: float) -> list[float]:
    return [value / temperature for value in logits]

def top_k_filter(logits: Sequence[float], k: int) -> list[float]:
    keep = {index for index, _value in sorted(enumerate(logits), key=lambda item: item[1], reverse=True)}
    return [value if index in keep else -1.0e9 for index, value in enumerate(logits)]

def sample_token(probabilities: Sequence[float], rng: random.Random) -> int:
    threshold = rng.random()
    cumulative = 0.0
    for index, probability in enumerate(probabilities):
        cumulative += probability
        if threshold <= cumulative:
            return index
    return len(probabilities) - 1

def generate(
    params: GPTParams,
    config: GPTConfig,
    start_ids: Sequence[int],
    max_new_tokens: int,
    temperature: float,
    top_k: int,
    rng: random.Random,
) -> list[int]:
    tokens = list(start_ids)
```

```

generated = []
for _ in range(max_new_tokens):
    logits = gpt_forward(tokens, params, config)
    probabilities = softmax(top_k_filter(temperature_scale(logits, temperature), top_k))
    next_token = sample_token(probabilities, rng)
    tokens.append(next_token)
    generated.append(next_token)
return generated

```

Run with `python3 src/python/chapter_demos.py`.

10.6 Weight Tying in Detail

Why does weight tying work so well?

- The embedding $E[i]$ is learned to represent token i such that tokens that appear in similar contexts have similar embeddings.
- The unembedding $\text{logits}[i] = h \cdot E[i]$ measures the alignment between the model’s prediction vector h and token i ’s embedding.
- If h is pointing in the direction of “tokens that follow this context,” and $E[i]$ represents token i ’s meaning, then tokens that are semantically appropriate will naturally score higher.

Weight tying forces the model to use a single geometric space for both “what a token means” and “how to predict a token” — an elegant constraint that regularizes learning.

10.7 Key Takeaways

- The *unembedding* projects the final hidden state to logits: $\text{logits} = hW_u \in \mathbb{R}^{|V|}$.
- *Weight tying* ($W_u = E^\top$) reuses the embedding matrix and reduces parameters.
- *Softmax* converts logits to a probability distribution.
- *Temperature* controls sampling sharpness; *top-K* restricts the candidate set.
- *Autoregressive generation*: sample one token at a time, feeding each back as input.

10.8 The Complete Picture

You have now seen every piece:

```

Input text
  ↓ tokenizer
Token IDs: [3, 7, 12, 5, ...]
  ↓ embedding matrix E
Token embeddings:  $[T \times d]$ 
  ↓ + positional encoding
 $X = X + PE: [T \times d]$ 
  ↓ transformer block  $\times N$ 
    LayerNorm
    Multi-Head Attention (Q, K, V projections + causal mask + softmax + weighted sum)
    Residual
    LayerNorm
    Feed-Forward Network (expand  $\rightarrow$  GELU  $\rightarrow$  contract)
    Residual
  ↓
 $X_{\text{final}}: [T \times d]$ 
  ↓ final LayerNorm + unembedding
Logits:  $[|V|]$ 
  ↓ softmax (+ temperature + top-K)
 $P(\text{next token}): [|V|] \rightarrow \text{sample} \rightarrow \text{next token ID}$ 

```

Every box in that diagram corresponds to one chapter of this book.

11 Loss — How the Model Learns

So far we have traced the forward pass from end to end. But none of that knowledge was conjured from thin air. The model *learned* by making predictions, measuring how wrong they were, and nudging every weight in the right direction — billions of times.

This chapter is about that measurement: the *loss function*.

11.1 The Idea

The model reads a sentence and, at each position, guesses what word comes next. Those guesses are never perfect — but how do we measure *how wrong* a guess is?

We need a single number that captures “wrongness” in a way that is small when the model is doing well and large when it is doing badly. That number is the *loss*. And crucially, it needs to be computable from the model’s output so we can use it to improve the model.

Here is the key insight: the model already outputs a probability for every word in the vocabulary. If the true next word is “cat” and the model assigns a 90% probability to “cat,” that is a good prediction — loss should be low. If the model assigns only 1% probability to “cat” (spreading the rest across unrelated words), that is a bad prediction — loss should be high.

The *cross-entropy loss* formalizes this: it measures how much probability the model assigned to the *correct* answer. A perfect model assigns 100% to the right word — loss is zero. A confused model spreads probability thinly — loss is high.

11.2 The Problem: Evaluating a Probability Distribution

After the forward pass, the model produces a probability distribution over the vocabulary:

$$P(\cdot \mid t_1, \dots, t_t) \in \mathbb{R}^{|V|}$$

We know the true next token t_{t+1} (because it comes from the training text). The question is: *how do we turn “how wrong is this distribution” into a single differentiable number?*

The answer is *cross-entropy loss*.

11.3 Cross-Entropy Loss

11.3.1 The Formula

For a single prediction at position t :

$$\mathcal{L}(t) = -\log P(t_{t+1} \mid t_1, \dots, t_t)$$

That's it. Negative log of the probability assigned to the *true* next token.

11.3.2 Why Negative Log?

The model outputs a probability $p \in (0, 1]$. We want a loss that is zero when the prediction is perfect ($p = 1$), large when the model is wrong ($p \rightarrow 0$), and differentiable everywhere.

$-\log(p)$ delivers all three:

Predicted probability p	$-\log(p)$	Interpretation
1.00	0.00	Perfect
0.80	0.22	Pretty good
0.50	0.69	Random guess
0.10	2.30	Mostly wrong
0.01	4.61	Very wrong
0.001	6.91	Catastrophically wrong

Suppose the vocabulary has four tokens and the model assigns probabilities $[0.1, 0.6, 0.2, 0.1]$. If the correct token index is 1, then the probability of the correct answer is 0.6. The loss is:

$$\mathcal{L} = -\log(0.6) \approx 0.511$$

If the model assigned probability 1.0 to the correct token, the loss would be 0.0. If it assigned 0.01, the loss would jump to 4.605. Low confidence in the correct answer is expensive, as Figure 11.1 shows.

Note

Math Minute — Logarithms

$\log(p)$ for $p \in (0, 1]$ is always ≤ 0 . It passes through $(1, 0)$ and tends to $-\infty$ as $p \rightarrow 0$. Because we negate it, $-\log(p) \geq 0$: zero means perfect, large means wrong. All modern ML uses natural log (base e), so loss is measured in *nats*. Using base 2 gives *bits*.




Figure 11.1: Cross-entropy loss grows steeply as the model assigns less probability to the correct token

11.3.3 The Full Training Loss

A single document of length T contributes:

$$\mathcal{L}_{\text{doc}} = \frac{1}{T} \sum_{t=1}^T -\log P(t_{t+1} | t_1, \dots, t_t)$$

11.3.4 Perplexity

Loss as nats is hard to interpret intuitively. *Perplexity* converts it to something more concrete:

$$\text{PPL} = e^{\mathcal{L}}$$

Perplexity is the *effective branching factor*: on average, how many equally likely choices does the model think there are at each step?

Loss $\mathcal{L}$	Perplexity	Interpretation
0.00	1.0	Perfect — only one plausible token
0.69	2.0	Two equally likely tokens
2.30	10.0	Ten equally likely tokens
4.61	100	Random guess over 100 tokens
6.91	1000	Very confused

GPT-2 (1.5B) reached ~18 perplexity on WikiText-103. GPT-4 is estimated well below 5 on standard benchmarks. Early training starts around 100–1000; loss curves falling to perplexity ~10–20 signals the model has learned real language structure.

i Note

Math Minute — Why Exponential?

If a model assigns uniform probability $1/k$ to each of k options, the cross-entropy is $-\log(1/k) = \log k$. Exponentiating: $e^{\log k} = k$. So perplexity of k means the model behaves like a uniform distribution over k choices. Perplexity 1 = model is certain. Perplexity $|V| = 50,000$ = model knows nothing.

11.3.5 Teacher Forcing

Notice the formula conditions on *true* past tokens, not the model's *own predictions*. This is called *teacher forcing*: during training, we always feed the ground-truth context.

This makes training stable, parallelizable, and simple. The causal mask from Chapter 4 is what enables this — position t can only attend to positions $1 \dots t$, so the model's prediction at position t is causally correct even when processing all positions in parallel.

11.3.6 Cross-Entropy as Information Theory

Cross-entropy has deep roots. The *cross-entropy* between true q and predicted p is:

$$H(q, p) = - \sum_j q_j \log p_j$$

When q is one-hot, only the true token's term survives: $H(q, p) = -\log p_{t_{+1}}$. Exactly our loss.

Minimizing cross-entropy is equivalent to minimizing the KL divergence between the model's predictions and the true data distribution. This is why cross-entropy is the natural loss for language models.

11.4 The Matrix: Worked Example

Let's trace the loss for a 4-token sequence ["The", "cat", "sat", "on"] with $|V| = 5$ (toy vocabulary).

11.4.1 Forward Pass Outputs

The model produces logits at each position. After softmax:

```
Position 0 (predicting token after "The"):  
  P = [0.05, 0.10, 0.60, 0.20, 0.05] ← true next = token 3  
  P(true) = 0.20, loss = -log(0.20) = 1.609  
  
Position 1 (predicting after "The cat"):  
  P = [0.10, 0.05, 0.15, 0.10, 0.60] ← true next = token 4  
  P(true) = 0.60, loss = -log(0.60) = 0.511  
  
Position 2 (predicting after "The cat sat"):  
  P = [0.40, 0.20, 0.15, 0.15, 0.10] ← true next = token 0  
  P(true) = 0.40, loss = -log(0.40) = 0.916  
  
Position 3 (predicting after "The cat sat on"):  
  P = [0.05, 0.05, 0.80, 0.05, 0.05] ← true next = token 2  
  P(true) = 0.80, loss = -log(0.80) = 0.223
```

11.4.2 Total Loss

$$\mathcal{L} = \frac{1}{4}(1.609 + 0.511 + 0.916 + 0.223) = 0.815$$

$$\text{PPL} = e^{0.815} \approx 2.26$$

The model is effectively choosing between about 2.3 equally likely tokens on average.

Figure [Figure 11.2](#) shows cross-entropy loss as softmax probability followed by the negative log of the true-token probability.

11.5 Python Implementation

`cross_entropy_loss` computes $-\log P(\text{true token})$ for a single position.

```
def cross_entropy_loss(logits: Sequence[float], true_id: int) -> float:  
    probabilities = softmax(logits)  
    return -math.log(max(probabilities[true_id], 1.0e-12))
```




Figure 11.2: Cross-entropy loss

`sequence_loss` averages the per-position losses into one scalar — the training signal for a complete sequence. `perplexity` converts the mean loss back to an interpretable scale.

```
def sequence_loss(logits_list: Sequence[Sequence[float]], true_ids: Sequence[int]) -> float:
    losses = [cross_entropy_loss(logits, true_id) for logits, true_id in zip(logits_list, true_ids)]
    return sum(losses) / len(losses)

# tag::perplexity[]
def perplexity(loss: float) -> float:
    return math.exp(loss)
# end::perplexity[]
```

The softmax-cross-entropy gradient has a famously clean closed form: $\partial \mathcal{L} / \partial z_j = P(j) - \mathbf{1}[j = \text{true}]$.

```
def softmax_cross_entropy_grad(logits: Sequence[float], true_id: int) -> Vector:
    grad = softmax(logits)
    grad[true_id] -= 1.0
    return grad
```

Run with `python3 src/python/chapter_demos.py`. Expected output:

Per-position losses:

```
pos 0 (true=3): loss=1.609 P(true)=0.200
pos 1 (true=4): loss=0.511 P(true)=0.600
pos 2 (true=0): loss=0.916 P(true)=0.400
pos 3 (true=2): loss=0.223 P(true)=0.800
```

Mean loss: 0.815

Perplexity: 2.26

11.6 Key Takeaways

- *Cross-entropy loss* is $-\log P(\text{true next token})$. Zero when perfect, unbounded when wrong.
- *Teacher forcing* feeds ground-truth context at training time, enabling full parallelism via causal masking.
- *Perplexity* $= e^{\mathcal{L}}$ is the effective branching factor — more intuitive than raw nats.
- *The softmax-CE gradient* is $P(j) - \mathbf{1}[j = t_{t+1}]$: a closed-form, numerically stable backprop step.
- *Cross-entropy* = *KL divergence* (since labels are one-hot), so training minimizes the KL from model predictions to the true data distribution.
- The loss is the *single* signal that shapes every weight: embeddings, attention projections, FFN weights, layer norms — all updated from this one number per token.

i Note

What's next? Chapter 9 gave us a loss: a single number measuring how wrong the model is. But knowing *how wrong* we are is only half the battle. We also need to know *which weights* are responsible for the error, and *by how much to change each one*. That is the job of *training* — Chapter 10.

12 Training — Teaching the Model

Chapter 9 gave us a loss: a single number measuring how wrong the model is. But knowing *how wrong* we are is only half the battle. We also need to know *which weights* are responsible for the error, and *by how much to change each one*.

That is the job of *training*.

12.1 The Idea

The loss tells us *how wrong* the model was at its last prediction. Training uses that wrongness to *improve* the model — so it does better next time.

Here is the core loop, in plain terms:

1. *Forward pass*: feed a sentence into the model, let it predict the next word at every position, measure how wrong those predictions were (the loss).
2. *Backward pass*: trace back through every calculation the model just made and figure out — for each weight — whether increasing it would have made the loss higher or lower, and by how much. This produces a *gradient* for every weight.
3. *Update*: nudge each weight slightly in the direction that reduces the loss. Weights that caused big errors get bigger nudges; weights that barely mattered get smaller nudges.
4. *Repeat*: billions of times, across trillions of words.

That is the entire training algorithm.

12.2 The Goal: Move the Loss Down

The model has millions of parameters. Call them collectively θ . The loss $\mathcal{L}(\theta)$ is a function of all of them. We want to find θ that minimizes the loss.

The strategy is *gradient descent*: repeatedly take a small step downhill on the loss surface.

$$\theta \leftarrow \theta - \eta \cdot \nabla_{\theta} \mathcal{L}$$

$\eta > 0$ is the *learning rate* — how big a step to take. $\nabla_{\theta}\mathcal{L}$ is the *gradient*: a collection of partial derivatives, one per parameter.

i Note

Math Minute — Partial Derivatives

A partial derivative $\frac{\partial \mathcal{L}}{\partial w}$ answers: *if we increase w by a tiny ϵ while holding every other weight fixed, by how much does $\mathcal{L}$ change?*

A large positive value means w is pushing loss up — decrease it. A negative value means the opposite. Zero means w is momentarily irrelevant to the loss.

Notation: ∂ (curly d) instead of d signals “partial” — other variables are held constant.

12.3 The Chain Rule

Backpropagation is the chain rule from calculus, applied to a computation graph.

Suppose the loss depends on weight w through a chain of intermediate values:

$$w \xrightarrow{f} z \xrightarrow{g} \mathcal{L}$$

The chain rule says:

$$\frac{d\mathcal{L}}{dw} = \frac{d\mathcal{L}}{dz} \cdot \frac{dz}{dw}$$

If we know how $\mathcal{L}$ varies with z , and how z varies with w , we multiply to get how $\mathcal{L}$ varies with w .

In a transformer, the chain has hundreds of layers. Backprop starts at the loss and works backward, multiplying local derivatives as it goes. Each layer only needs two things: the gradient arriving from the layer above (called δ , “delta”), and its own local derivative.

12.3.1 Step 1: A gradient update for one weight

```
def sgd_update_scalar(weight: float, gradient: float, learning_rate: float) -> float:
    return weight - learning_rate * gradient
```

If the gradient is positive (loss rises when w increases), we subtract — decreasing w . If negative, we add. The learning rate η scales how large the step is.

This is the entire idea. Everything that follows is this rule, applied to millions of weights simultaneously.

12.4 Backprop Through Softmax and Cross-Entropy

Chapter 9 stated the softmax-CE gradient without proof:

$$\frac{\partial \mathcal{L}}{\partial z_j} = P(j) - \mathbb{1}[j = t]$$

where z_j are the logits and t is the true token.

For the true token $j = t$: gradient is $P_t - 1$ (negative, so we push z_t up). For every other token: gradient is P_j (positive, so we push those logits down). The model is nudged to be more confident about the correct answer.

```
def softmax_cross_entropy_grad(logits: Sequence[float], true_id: int) -> Vector:
    grad = softmax(logits)
    grad[true_id] -= 1.0
    return grad
```

12.5 Backprop Through a Linear Layer

The most common operation in a transformer is a *linear layer*: $y = Wx$.

Given the upstream gradient $\delta = \frac{\partial \mathcal{L}}{\partial y}$ (arriving from the layer above), the chain rule gives:

$$\frac{\partial \mathcal{L}}{\partial W} = \delta \cdot x^\top \quad \frac{\partial \mathcal{L}}{\partial x} = W^\top \cdot \delta$$

The first equation tells us how to update W . The second passes the gradient backward to whatever fed x into this layer.

12.5.1 Step 2: Linear layer backward pass

```
def linear_backward(delta: Vector, weights: Matrix, x: Vector) -> tuple[Matrix, Vector]:
    grad_w = [[delta_i * x_j for x_j in x] for delta_i in delta]
    grad_x = [sum(weights[i][j] * delta[i] for i in range(len(delta))) for j in range(len(x))]
    return grad_w, grad_x
```

12.6 Accumulating Gradients Across the Sequence

We compute the softmax-CE gradient independently at each position t . But all positions share the same unembedding matrix W_u (and the same transformer weights). Their gradient contributions must be *summed* before updating any weight.

12.6.1 Step 4: Summing gradient contributions

```
def accumulate_gradients(gradients: Sequence[Matrix]) -> Matrix:
    rows, cols = shape(gradients[0])
    total = make_matrix(rows, cols)
    for gradient in gradients:
        total = matrix_add(total, gradient)
    return total
```

The more a weight influenced the output at many positions, the larger (and more reliable) its accumulated gradient.

12.7 The Gradient Descent Step

With gradients accumulated for every parameter, the update rule is:

$$W \leftarrow W - \eta \cdot \frac{\partial \mathcal{L}}{\partial W}$$

12.7.1 Step 5: Updating a weight matrix in place

```
def sgd_update_matrix(weights: Matrix, gradient: Matrix, learning_rate: float) -> None:
    for i, row in enumerate(weights):
        for j, value in enumerate(row):
            row[j] = value - learning_rate * gradient[i][j]
```

Every entry W_{ij} moves a small step in the direction that reduces the loss.

12.8 One Training Step

Now we can assemble the full cycle: forward pass to get predictions and loss, backward pass to get gradients, update pass to improve every weight.

12.9 Watching the Loss Fall

Running thousands of training steps, the loss curve looks roughly like this:

Step	Loss	Perplexity	Interpretation
0	10.82	50,000	Random — model knows nothing
100	6.91	1,000	Ruling out most tokens
1,000	4.61	100	Learned basic frequency
10,000	2.30	10	Has rough grammar
100,000	0.92	2.5	Strong language model

The curve drops steeply at first (obvious mistakes are easy to fix) then more slowly and noisily (subtle patterns are harder and data points disagree).

Note

Why does loss get noisy later?

Early on, large gradients correct glaring errors. Later, gradients are smaller and point in slightly different directions for different training examples. The noise is not a bug: random variation in gradient direction helps the model escape poor local minima. This is the *stochastic* in *stochastic gradient descent* (SGD).

Figure Figure 12.1 shows gradients flowing backward from the loss through each layer while weights are updated.

12.10 Key Takeaways

- *Gradient descent* moves every weight downhill: $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$.
- *Partial derivatives* measure the sensitivity of the loss to each individual weight, holding all others fixed.




Figure 12.1: Backpropagation and weight updates

- *Backpropagation* is the chain rule applied layer by layer from loss to inputs. Each layer only needs the upstream gradient δ and its own local derivative — no global information required.
- *Linear layer backward*: $\frac{\partial \mathcal{L}}{\partial W} = \delta x^\top$, $\frac{\partial \mathcal{L}}{\partial x} = W^\top \delta$. The same two matrices as the forward pass, just transposed and multiplied in a different order.
- *Softmax-CE gradient* is $P(j) - \mathbb{1}[j = t]$: the model's predicted probability minus the true label. Nearly zero when correct and confident; large when wrong.
- *Shared weights accumulate gradients* from every position in the sequence before any update is applied.
- The six-step cycle — forward, loss, backward, accumulate, update, repeat — is all of training. Every weight in the model, over billions of tokens, updated by exactly this loop.

Part V

Modern Extensions

This part connects the baseline transformer to modern GPT implementations. It keeps the simple model from earlier chapters as the reference point, then shows which production refinements change speed, memory use, context length, or architecture.

- In Chapter [13](#), you will see rotary position embeddings, grouped-query attention, FlashAttention, and major architectural variants.

13 Modern GPT

The previous chapters built a complete GPT from scratch. Real production models add several refinements that improve efficiency, context length, and capability. This chapter surveys the most important innovations beyond the baseline transformer.

13.1 Rotary Positional Encoding (RoPE)

The sinusoidal positional encoding in Chapter 5 adds a fixed signal to each embedding *before* the attention layers. *RoPE* (Su et al., 2021) takes a different approach: it encodes position directly into the query and key vectors inside each attention layer by rotating them.

For a query vector q at position m and key vector k at position n :

$$\langle R_m q, R_n k \rangle = \langle q, R_{n-m} k \rangle$$

where R_θ is a block-diagonal rotation matrix. The dot product depends only on the *relative* position $n - m$, not the absolute positions. This gives the model a natural sense of distance between tokens.

Key advantages of RoPE:

- Generalises to sequences longer than those seen during training.
- No extra parameters — the rotation is deterministic.
- Compatible with linear attention variants.
- Used in LLaMA, Mistral, Qwen, and most open-weight models.

In 2D, a rotation by θ looks like:

$$R_\theta = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

For a d -dimensional vector, $d/2$ independent 2D rotations are applied to pairs of dimensions.




Figure 13.1: RoPE rotates query and key vectors by position — their dot product captures relative distance

13.2 Multi-Query and Grouped-Query Attention

Standard multi-head attention (Chapter 7) creates separate Q, K, V projections for every head. The key and value matrices dominate GPU memory during inference.

Multi-Query Attention (MQA) (Shazeer, 2019) uses a single K and V shared across all heads, reducing the KV cache by a factor of H (number of heads).

Grouped-Query Attention (GQA) (Ainslie et al., 2023) is a middle ground: G groups of heads share a single K, V , where $1 \leq G \leq H$. Setting $G = 1$ recovers MQA; setting $G = H$ recovers standard MHA.

Method	KV heads	Memory	Quality
MHA	H	1x	Baseline
GQA	H/G	1/G x	Near-MHA
MQA	1	1/H x	Slight drop

GQA is used in LLaMA 3, Mistral, and Gemma. It allows inference to fit in less memory without significant quality degradation.

13.3 Flash Attention

Standard self-attention computes the full $n \times n$ attention matrix explicitly:

$$\text{Attn} = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V$$

For a sequence of length n , this requires $O(n^2)$ memory — prohibitive for long contexts.

Flash Attention (Dao et al., 2022) reorders the computation using *tiling*: it processes blocks of the sequence in GPU SRAM (fast on-chip memory) and never materialises the full $n \times n$ matrix in HBM (slow off-chip memory).

The result is mathematically identical to standard attention, but:

- Memory: $O(n)$ instead of $O(n^2)$.
- Speed: 2–4× faster than PyTorch’s built-in attention in practice.
- Enables context windows of 32k–128k tokens at practical batch sizes.

Flash Attention 2 and 3 added further improvements (better parallelism, support for GQA). It is now the default in all major frameworks.

Note

Flash Attention is a hardware-aware algorithm, not a new mathematical operation. The inputs and outputs are identical to standard attention.

13.4 Alternative Architectures

The transformer is not the only architecture for sequence modelling. Several alternatives challenge or complement it.

13.4.1 Mamba (State Space Models)

Mamba (Gu & Dao, 2023) is based on *selective state space models* (SSMs). Instead of attending to all past tokens, it maintains a compressed hidden state that is updated recurrently — like an RNN, but with a selective mechanism that decides what to remember.

Key properties:

- Linear time in sequence length ($O(n)$ vs $O(n^2)$ for attention).
- No attention matrix — the context is compressed into a fixed-size state.
- Competitive with transformers on language modelling benchmarks.
- Less effective at in-context retrieval tasks (where attention excels).

13.4.2 Diffusion Language Models (DiffusionGemma)

Autoregressive models like GPT generate text *left-to-right*, one token at a time. *Diffusion language models* (DLMs) generate all tokens in parallel by iterative denoising.

MDLM and *Gemma-based diffusion models* start from a fully masked sequence and refine it over T denoising steps. Each step predicts a cleaner version of the entire sequence.

Advantages: - Can revise earlier tokens after seeing later context. - Faster generation at inference time (fewer sequential steps than token length). - Emerging area — quality trails autoregressive models as of 2025.

13.4.3 Mixture of Experts (MoE)

Mixture of Experts replaces the single feed-forward network in each transformer block with E expert networks and a *router* that activates only K of them per token.

- Mistral 8x7B and GPT-4 are believed to use MoE.
- Total parameters can be large while compute per token stays small.
- Trade-off: routing instability, expert load imbalance.

13.5 Key Takeaways

Innovation	Benefit
RoPE	Relative position encoding, better length generalisation
GQA / MQA	Smaller KV cache, faster and cheaper inference
Flash Attention	Linear memory, 2-4x faster attention, long context
Mamba (SSM)	Linear-time sequence modelling, no attention matrix
Diffusion LMs	Parallel generation, bidirectional revision
MoE	Larger model capacity without proportional compute cost

The field moves fast. Each of these innovations addresses a concrete bottleneck — memory, speed, context length, or capability — that became critical as models scaled. Understanding the baseline GPT from Chapters 3–12 makes every one of these extensions legible.

A microGPT in Python — Complete Runnable Code

This appendix assembles every Python snippet from the book into three files under `src/python/`, with all dependencies resolved. The implementation uses only the Python standard library.

A.1 Files

- `common.py` — forward-pass building blocks (embeddings, attention, FFN, transformer, GPT forward)
- `train.py` — loss functions and gradient descent (cross-entropy, backprop, SGD)
- `inference.py` — generation and end-to-end demo

A.2 Installation & Running

```
# Run the full demo
python3 src/python/inference.py

# Run every chapter's end-to-end code check
python3 src/python/run_book_code.py
```

A.3 End-to-End Demo

```
def demo(seed: int = 7) -> dict[str, object]:
    config = GPTConfig(vocab_size=100, d_model=32, num_heads=4, num_layers=2, max_seq_len=64)
    model_rng = random.Random(seed)
    sample_rng = random.Random(seed + 1)
    params = make_gpt_params(config, model_rng)
    logits = gpt_forward([1, 2, 3, 4, 5], params, config)
```

```

probabilities = softmax(logits)
top_token = max(range(len(probabilities)), key=probabilities.__getitem__)
generated = generate(params, config, [1, 2, 3], 10, temperature=0.8, top_k=10, rng=sampler)
return {
    "config": config,
    "parameters": count_parameters(config),
    "logits_first_10": logits[:10],
    "top_token": top_token,
    "top_probability": probabilities[top_token],
    "generated": generated,
    "full_sequence": [1, 2, 3] + generated,
}

```

A.4 microGPT Architecture Summary (Reference)

Component	Input	Output	Parameters
Token Embedding	[T] ints	$[T \times d]$	$ V \times d$
Positional Embedding	[T] positions	$[T \times d]$	$T_{max} \times d$
$\times N$ Transformer Blocks			
LayerNorm 1	$[T \times d]$	$[T \times d]$	$2d$
Multi-Head Attn	$[T \times d]$	$[T \times d]$	$4d^2$
Residual	$[T \times d]$	$[T \times d]$	0
LayerNorm 2	$[T \times d]$	$[T \times d]$	$2d$
FFN	$[T \times d]$	$[T \times d]$	$8d^2$
Residual	$[T \times d]$	$[T \times d]$	0
Final LayerNorm	$[T \times d]$	$[T \times d]$	$2d$
Unembedding	[d]	$[V]$	$d \times V $ (tied)

Total parameters: $2|V|d + T_{max} \cdot d + N(12d^2 + 4d) + 2d$

For GPT-2 small: $d=768$, $N=12$, $|V|=50257$, $T_{max}=1024 \rightarrow \sim 117M$ params.